



Bangladesh University of Engineering and Technology

Department of Computer Science and Engineering

Kinetic-KAN

Solver-Aware Neural Dynamics

Kolmogorov–Arnold networks as learned vector fields for neural ordinary differential equations: how integrator order, basis design, stiffness and training budget shape what the model learns

RUNGE–KUTTA INTEGRATION ♦ BASIS-FUNCTION DESIGN ♦ NUMERICAL STABILITY ♦ EXTRAPOLATION

Project Final Report

SUBMITTED BY

2105032 Nawriz Ahmed Turjo
2105033 Abhishek Roy
2105043 Monjur Hossain Khan (Shovon)
2105048 Shams Hossain Simanto
2105055 Md. Abrar Jahin

SUBMITTED TO

Niaz Rahman
Lecturer
Department of CSE, BUET

October 2, 2026

Contents

1	Introduction	7
1.1	Background: Kolmogorov-Arnold Networks as vector fields	7
1.2	Architecture, and what was built rather than imported	7
1.3	What the base paper fixed, and what this project varies	9
1.4	Why training itself needs to be treated as an object of study	9
1.5	Structure of this report	10
2	Baseline Recipe, Evaluation Protocol and Reproduction	10
2.1	Training recipe	11
2.2	Evaluation protocol	11
2.3	Computational budget: why 10,000 epochs	12
2.4	Baseline reproduction: KAN-ODE against parameter-matched MLP-ODEs	12
3	Benchmark Ablation and Cross-Domain Stability Analysis	15
3.1	The six integrators, and a check that each has its design order	15
3.2	Solver ablation	18
3.3	Basis functions: definitions, conditioning and locality	19
3.4	Sensitivity to step size and observation noise	21
3.5	What the trained vector fields actually learned	22
3.6	Training dynamics across the sweeps	25
3.7	Diagnosing and fixing two cross-domain failures	28
3.8	Extended extrapolation: the KAN advantage widens with horizon	30
3.9	Pendulum energy dissipation: an independent physical check	30
4	Adjoint Sensitivity vs. Direct Autograd: Memory and Wall-Clock Cost	31
5	Gradient-Norm Dynamics and Solver Order	33
6	Learnable Basis-Function Blending	35
7	Stiffness–Solver Stability and Loss-Landscape Analysis	38
7.1	The stability question resolves quickly	38
7.2	What actually drives the failures: initialization	38
7.3	The mechanism: a genuine spurious equilibrium	39
7.4	Weight-space loss landscape: an independent measurement	39
8	Sparse Regression Comparison and Outbreak-Curve Validation	43
8.1	Noise robustness: KAN-ODE vs. sparse symbolic regression	43
8.2	Edge pruning: testing whether a term-by-term comparison is even possible	44
8.3	Testing whether SIR’s fixes transfer to a non-mechanistic epidemic curve	45
9	Validation: Epoch-Budget Sensitivity and Implementation Audit	47
9.1	Design	47
9.2	Results	47
9.3	Implementation audit against the base paper	50
10	Discussion	51
11	Conclusion	52

List of Figures

- 1 The KAN-ODE training loop as implemented. The integrator calls the network once per Runge-Kutta stage, so one training epoch on $[0, 3.5]$ evaluates $\mathbf{f}_\theta$ 70s times (35 sampling intervals, two sub-steps each, s stages per step). The loss gradient is obtained by differentiating through every one of those evaluations. Numbered badges mark the three components this report varies: the edge basis (Section 3), the integrator (Sections 3, 5, 7) and the gradient method (Section 4). 8
- 2 Tensor flow through one KDense layer (Equation 3). Only the basis-expansion step differs between the seven basis families studied in Section 3; every basis function has the same signature, $(\mathbf{U}, \mathbf{z}, h) \mapsto \Psi \in \mathbb{R}^{B \times d_{\text{in}} \times G}$, so swapping the basis changes one tensor and nothing else. 8
- 3 Time axis of the Lotka-Volterra protocol. Tick marks above the bar are multiples of the true period. The far-horizon window $(14, 28]$ is used only for the stress tests of Sections 3 and 9; no model is retrained or reselected for it. 11
- 4 Training MSE per epoch for KAN-ODE and the paper’s parameter-matched tanh MLP-ODE, trained with the paper’s own learning rate and initialization (fourth column of Table 3). KAN-ODE descends steadily and reaches 10^{-4} at epoch 9617; the tanh MLP reaches 10^{-2} at epoch 3461 but never 10^{-3} within 10^4 epochs. 13
- 5 Predicted against true populations out to $t = 28$ for the three models, with no retraining. KAN-ODE and SiLU-MLP are visually indistinguishable from the truth. The paper’s tanh MLP under its own learning rate and initialization (bottom, fourth column of Table 3) follows the training window and the first periods closely, then runs slightly slow and lags visibly by the far window $(14, 28]$ 13
- 6 Stage nodes $t_n + c_i h$ and solution weights b_i of the six implemented methods, drawn from the coefficients in the code. Heun and Midpoint use the same number of evaluations but sample at opposite places: Heun averages the two ends of the step, Midpoint uses only a slope at its centre. Tsit5 reaches order five with large weights of both signs ($b_5 = -3.29$), a cancellation that is harmless in exact arithmetic and is what buys its small error constant. 16
- 7 Global error at $t = 3.5$ of each implemented integrator applied to the true Lotka-Volterra field. (a) Against step size, with fitted slopes; the dashed line is the training step $h = 0.05$. (b) Against the number of field evaluations, the quantity that sets training cost. At $h = 0.4$ both second-order methods blow up, and the fifth-order methods reach the round-off floor near 10^{-13} 17
- 8 (a) Stability boundaries $|R(z)| = 1$ computed from each implemented tableau, with $h\lambda$ for the true Jacobian (dots) and the learned Jacobian of the Tsit5-trained KAN (crosses) at every state on the orbit, $h = 0.05$; the inset enlarges the neighbourhood of the origin. (b) Per-step amplitude error for an undamped oscillation; the dotted line marks $h\omega_0$ at the training step. 17
- 9 Extrapolation MSE of each trained model (rows) integrated with each solver (columns), $h = 0.05$. The diagonal (bold) is the configuration it was trained in. 19
- 10 Training wall-clock against extrapolation MSE for the solver and basis sweeps. The dashed step is the Pareto front: Euler, Midpoint, Tsit5 and B-spline. Midpoint dominates RK4 and DOPRI5 outright. 19
- 11 Phase portraits of the Euler-, Midpoint- and Tsit5-trained RBF models on $[0, 14]$ (all three plots are titled “KAN-ODE (RBF)” because the basis is shared). Each model is integrated with its own training solver. All three stay on the true orbit to plotting accuracy, including Euler, whose $28\times$ larger extrapolation MSE is phase error along the orbit rather than a visible drift off it. 19
- 12 The seven implemented basis families on the $G = 5$ grid (vertical lines are the knots z_g), with the $g = 4$ function highlighted and the 2-norm condition number of the design matrix in each title. The blue band is the interval $u \in [0.25, 1]$ that first-layer inputs actually reach on the Lotka-Volterra orbit. The last panel shows κ_2 against grid size; dashed and dotted grey lines are IQF and RSWAF. 20
- 13 Locality of one coefficient update, $c_2 \leftarrow c_2 + \delta$, for three basis families (uniform-knot cubic spline shown for clarity). Since ϕ is linear in $\mathbf{c}$, the change is exactly $\delta \psi_2$ (shaded). Compact support confines it to its knot span; the RBF’s Gaussian tail still reaches the data interval (blue); a global polynomial changes the edge everywhere. 21

14	Phase portraits on $[0, 14]$ for four of the seven bases (Tsit5). B-spline, RBF and Chebyshev all trace the true orbit closely at this scale; Chebyshev's $5\times$ larger extrapolation error is phase error rather than a visible departure. The Newton-basis model visibly misses the orbit near its upper and right-hand extremes.	21
15	Left: error against sampling interval, with $O(\Delta t)$ and $O(\Delta t^2)$ reference slopes; changing Δt also changes the number of training samples. Right: error against observation-noise level, with least-squares power-law fits; the dotted line is the noise variance σ^2	22
16	Trajectories of each learned field from four initial conditions (coloured dots) over $t \in [0, 40]$, against the four corresponding true orbits (grey). Red crosses mark the learned equilibria, black plus signs the true one. Columns: KAN-ODE, SiLU MLP-ODE, and the paper's tanh MLP-ODE under the paper's own learning rate and initialization; rows: 10^4 and 5×10^4 epochs. Every learned field attracts all four onto the training orbit except the 10^4 -epoch tanh MLP, under which the trajectory from $(0.6, 0.6)$ diverges.	24
17	Field error $\ \mathbf{f}_\theta(\mathbf{u}) - \mathbf{f}(\mathbf{u})\ $, normalised by the mean true field magnitude on the orbit (4.82), clipped to $[10^{-3}, 1]$. Thick black: training segment; blue: rest of the orbit; plus sign: true equilibrium. Columns and rows as in Figure 16.	24
18	Relative drift of the Lotka-Volterra first integral H (Equation 5). (a) The exact field integrated by each solver at $h = 0.05$: pure truncation drift (Euler's trajectory leaves the domain of H at $t = 19.4$). (b) Each trained model integrated by its training solver. Shaded: training window; dotted: end of the scored horizon.	25
19	Pointwise trajectory error on $[0, 28]$. (a) Solver and basis choice at 10^4 epochs (Midpoint and Tsit5 overlap). (b) KAN-ODE and MLP-ODE at 10^4 and 5×10^4 epochs.	25
20	First-layer edge functions (spline plus residual branch) of the trained RBF and B-spline models, each input varied over its range on the orbit with the other held at its mean. The edge with the largest range is highlighted.	26
21	Test MSE during training for the six solvers (RBF basis), logged every ten epochs; the dotted line is 10^{-4}	26
22	Test MSE during training for the seven bases (Tsit5 solver), logged every ten epochs; the dotted line is 10^{-4}	26
23	Fixed-pendulum trajectory reconstruction (left) and training/extrapolation loss curves (right), using the extended-window recipe.	28
24	SIR with all three fixes: reconstruction of S and I over the full horizon (left; the plot's generic legend labels them prey and predator) tracks the true dynamics closely; training and extrapolation loss curves (right) decrease together without the gradient overflow that terminated the original run.	29
25	Fixed pendulum ($\mu = 0.5$, $t_{\text{train}} = 5$): phase portrait (left) and total mechanical energy $E(\theta, \omega)$ (right). Inside the training window the predicted energy follows the true curve closely; past the split at $t = 5$ it stays above the truth and rises on 40 of 200 predicted steps, the under-damping described in the text.	31
26	Direct autograd (top) keeps every intermediate state u_i in memory so the backward pass can reuse it, giving $O(N_t)$ memory. Adjoint sensitivity (bottom) instead solves the backward-time equation of Equation 18, recomputing $u(t)$ as needed rather than storing it, giving $O(1)$ memory regardless of N_t	31
27	Peak memory vs. trajectory length $N_t \in \{36, \dots, 3201\}$: direct autograd grows linearly (≈ 0.045 MB per output point); adjoint stays flat (0.7% increase over an $89\times$ length increase). . .	32
28	Gradient-norm traces, all six solvers: large sustained spike bursts (one to two orders of magnitude above the local median) dominate every trace, including Tsit5, independent of solver order.	34
29	The hybrid-basis edge of Equation 20: every edge evaluates both bases in parallel, and a shared softmax gate produces the scalar weights α, β that multiply each basis's output before they are summed into $\phi(x)$ – the choice of basis is a trained quantity, not a decision fixed at construction time.	35

30	Learned α/β gate evolution: Lotka-Volterra settles into a genuine partial blend (left); the pendulum gate reaches $\approx 99.5\%$ RBF by epoch 250 and ends at 98% (right); it settles well before the overfitting onset discussed below.	36
31	Left: the 2000-epoch probe run's gate trace, showing the slow drift toward RBF that motivated the gate-learning-rate fix. Right: the gate trace for the 3000-epoch pre-overfitting-optimum run discussed below, settling to $\beta \approx 0.996$	36
32	Hybrid basis on the pendulum. Top: the full 10,000-epoch run; training loss keeps falling while the extrapolation (monitor) loss bottoms out near epoch 2950 and climbs back to ≈ 0.4 (left), and the extrapolated orbit collapses into a small wrong loop (right). Bottom: the same recipe stopped at 3000 epochs; train and monitor loss fall together (left) and the extrapolation follows the decaying spiral (right).	37
33	Stability verdict heatmap, 24-cell single-seed probe sweep, $\Delta t = 0.05$	39
34	Learned vector-field landscape $ f_\theta(\theta, \omega) $, per solver's own trained model. Left ($\mu = 2.0$): every solver's true equilibrium (star) sits partway up a slope, while a separate, lower point nearby is where each trajectory actually settles. Right ($\mu = 1.0$): Euler's star sits high on a ridge, well apart from where its trajectory spirals down and rests; for the other three solvers, star and resting point coincide at the valley floor.	40
35	$\mu = 2.0$, all four solvers: predicted angle, angular velocity, and error over time, alongside each solver's phase portrait.	40
36	$\mu = 1.0$, all four solvers: the same panel layout, showing Euler's trajectory diverging from the true angle while the other three track it closely.	40
37	Euler, $\mu = 2.0$ (trapped): the training-loss surface around the trained weights, viewed under three independent random direction-pairs. The basin near the trained point (the marked star) is broad and shallow in all three views; the third panel's steep spike on the far side is the outlier direction discussed below, not the basin itself.	41
38	Euler, $\mu = 1.0$ (also trapped): the same broad, shallow basin shape recurs at a different damping value under the same solver.	41
39	Tsit5, $\mu = 1.0$ (converged): the basin around the trained point is visibly narrower and drops noticeably lower, in every one of the three direction-pairs, than either trapped case above. . . .	42
40	Mean rise per (solver, μ), 3-seed average with seed-to-seed range as error bars, colored by verdict.	42
41	Converged fraction (left) and median training MSE (right) over 3 seeds, all 4 solvers $\times$ 6 μ (72 cells). The two agree everywhere: low converged-fraction cells are exactly the cells with elevated median error.	43
42	Training and extrapolation MSE vs. noise level, both methods.	44
43	Reconstructed trajectories at $\sigma = 0$ and $\sigma = 0.1$: sparse regression loses amplitude under noise; KAN-ODE keeps amplitude but drifts in phase.	44
44	Full-horizon MSE vs. edge-pruning fraction, all four noise levels.	45
45	Time-rescaling sweep: accuracy and gradient-spike ratio vs. the rescaling factor.	46
46	Outbreak-curve fit: default split (runaway), the vanishing-field-gate arm, and the fifteen-day-later split control.	46
47	Training and test loss of the six extended runs, as saved by each run. Training loss is logged every epoch; the test loss (full horizon $[0, 14]$) every ten epochs, and it spikes by up to two orders of magnitude throughout, which is why the comparison in the text uses a 500-epoch rolling median rather than raw values. Panel (f) shows the identical $[2, 50, 2]$ tanh architecture of panel (e), retrained with only the learning rate and initialization scale changed to the values in the paper's own published code.	48

List of Tables

1	Shared training recipe for all Lotka-Volterra runs.	11
2	Measured per-epoch cost, local development CPU against cloud (Kaggle) CPU, each averaged over the full run, and the projected cost of the paper’s 10^5 -epoch budget.	12
3	KAN-ODE against parameter-matched MLP-ODEs at 10^4 epochs. Epoch counts are the first epoch at which the training loss falls below the threshold. The last three rows are computed from the saved checkpoints: field error is $\ \mathbf{f}_\theta - \mathbf{f}\ _2$ averaged over the true orbit on $(3.5, 14]$ and divided by the mean $\ \mathbf{f}\ _2 = 4.82$ there; the period is measured from zero-crossings of the predicted prey signal over $[0, 28]$; H drift is $\max_t H(\hat{\mathbf{u}}(t)) - H_0 /H_0$ on $[0, 28]$. The fourth column is the identical literal baseline of the third, retrained under the paper’s own learning rate and initialization scale (Section 9.3); its 5×10^4 -epoch continuation is in Table 23.	14
4	Controlled breakdown of why the paper’s literal MLP baseline architecture, $[2, 50, 2]$ with hyperbolic tangent (tanh) activation, fails to train at this project’s default learning rate and initialization scale. All four runs use 2000 epochs and Glorot initialization at gain 1.0; none varies the initialization scale (see the paragraph following this table).	14
5	Fixed-step Tsit5 tableau as implemented (first six stages, rounded to six significant figures; the code carries 16). Row i lists c_i and $a_{i1}, \dots, a_{i,i-1}$; the last row lists the solution weights b_i	15
6	Order verification of the implemented integrators. “Max residual” is the largest $ \mathbf{b}^\top \Phi(\tau) - 1/\gamma(\tau) $ over all trees up to the design order ($1 + 1 + 2 + 4 + 9 = 17$ conditions at order five); $A^{(p+1)}$ is the principal error norm; $\hat{p}$ is the slope fitted to the work-precision data for $h \leq 0.1$ above the round-off floor.	16
7	Solver ablation (fixed RBF basis, $h = 0.05$, 240 parameters, 10^4 epochs). The last column integrates the <i>true</i> field with the same solver and step and scores it identically: the solver’s own truncation error, with no learning involved.	18
8	Basis ablation (fixed Tsit5 solver, $G = 5$ grid points, 240 parameters, 10^4 epochs). κ_2 is the design-matrix condition number of Figure 12.	21
9	Step-size and observation-noise sweeps (Tsit5, RBF, 10^4 epochs). Noise is added only to the training window; all errors are measured against the clean truth.	22
10	Numerical analysis of the trained fields. $\mathbf{u}^*$: learned equilibrium (Newton-Raphson from $(3, 1.5)$); λ : eigenvalues of $J_\theta(\mathbf{u}^*)$; period error from zero-crossings of the predicted prey signal over $[0, 28]$; H drift is $\max_t H(\hat{\mathbf{u}}) - H_0 /H_0$ on $[0, 28]$. True values: $\mathbf{u}^* = (3, 1.5)$, $\lambda = \pm 2.121i$, $T = 3.3175$, drift 0.	23
11	First epoch at which the training MSE falls below each threshold (“–”: not reached within 10^4 epochs).	27
12	Pendulum fix: best configuration versus the original failure.	28
13	SIR before and after the fixes, at the full 10,000-epoch budget. “Rescale + project” applies time rescaling and the conservation projection, which assume nothing beyond Equation 6; “+ gate” adds the vanishing-field gate. Mass error is $\max_t S + I + R - 1 $	29
14	Every stability intervention tested, including the rejected ones.	30
15	Extrapolation MSE as the horizon doubles, same trained models.	30
16	Peak VRAM and wall-clock, direct autograd vs. adjoint sensitivity.	32
17	Gradient-noise measure η vs. solver order, six solvers, $n = 6$	33
18	Hybrid basis vs. pure bases, full 10,000-epoch budget.	35
19	Original 24-cell sweep (single seed 42), best training MSE and verdict.	38
20	Loss-landscape rise, single direction-pair (seed 0).	41
21	Extrapolation MSE vs. observation noise σ	43
22	Which of SIR’s three fixes transfer to a non-mechanistic curve.	45
23	The five originally planned configurations at 10^4 epochs (Section 3) and at the extended budget, plus a sixth added after auditing the paper’s own optimizer settings for its MLP baseline (Section 9.3). “Clip” is the fraction of optimizer steps on which gradient clipping at norm 1.0 engaged; $\ \nabla\ _{\text{med}}$ is the median pre-clip gradient norm over the extended run. The sixth configuration ran with no clipping by design (see the paragraph following this table), so those two columns are blank throughout.	47

24 Which 10^4 -epoch conclusions survive the extended budget. 49

25 Field-by-field comparison with the base paper [6]. 50

1 Introduction

1.1 Background: Kolmogorov-Arnold Networks as vector fields

A standard Neural ODE [3] approximates an unknown dynamical system $\dot{\mathbf{u}}(t) = \mathbf{F}(\mathbf{u}(t))$ by a learned vector field $\mathbf{f}_\theta(\mathbf{u}(t))$, parameterized by a multilayer perceptron: fixed nonlinear activations on each node, with all learnable structure in the linear weights, $\mathbf{y} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$. A Kolmogorov-Arnold Network (KAN) [8] inverts this: every *edge* carries its own learnable univariate function $\phi(x)$, and each node simply sums its incoming edges, $y_j = \sum_i \phi_{i,j}(x_i)$, following the Kolmogorov-Arnold representation theorem – any continuous multivariate function on a bounded domain can be written as a finite composition of continuous univariate functions and addition,

$$f(\mathbf{x}) = \sum_{q=1}^{2n+1} \Phi_q \left(\sum_{p=1}^n \phi_{q,p}(x_p) \right), \quad (1)$$

so that no fixed nonlinearity needs to be chosen in advance – the network learns what shape each univariate function should take. In practice each edge function is represented as a fixed smooth residual plus a learnable expansion in a chosen basis. In the form implemented for this project,

$$\phi(x) = w_b b(x) + \sum_{i=1}^G c_i B_i(\text{norm}(x)), \quad b(x) = \text{silu}(x) = \frac{x}{1 + e^{-x}}, \quad \text{norm}(x) = \tanh x, \quad (2)$$

where $B_1, \dots, B_G$ are G basis functions attached to a uniform grid of knots $z_g \in [-1, 1]$ with spacing $h = 2/(G - 1)$, c_i are trainable coefficients, and w_b weights the fixed residual. (The original pykan form carries $G + k$ spline functions and a separate scale w_s on the expansion; with one coefficient per basis function that scale is redundant, and the implementation drops it.) The choice of B_i is a design decision, not a mathematical necessity: a Gaussian radial basis function (RBF) grid, $B_i(u) = \exp(-(u - z_i)^2/2h^2)$, gives each edge smooth support that decays quickly but never reaches zero, while a degree- k B-spline built by the Cox-de Boor recursion gives each edge compact local support, exactly zero outside its own knot span. KAN-ODEs [6] put this construction inside a Neural ODE’s vector field: $\dot{\mathbf{u}} = \mathbf{f}_\theta(\mathbf{u})$ with every edge of $\mathbf{f}_\theta$ built from Equation 2, and the resulting trajectory obtained by numerically integrating that field forward in time.

1.2 Architecture, and what was built rather than imported

Every KAN-ODE trained for this report shares one architecture: a two-layer network with widths $[2, 10, 2]$, each edge carrying the expansion of Equation 2 over a grid of $G = 5$ basis functions, wrapped inside an explicit Runge-Kutta integrator that advances the state and hands the resulting trajectory to the loss. Figure 1 shows the full training loop and marks the three components this report varies.

1.2.1 One KAN layer as two matrix products

The network is two KDense layers. For a batch $\mathbf{X} \in \mathbb{R}^{B \times d_{\text{in}}}$, a layer evaluates every edge function of Equation 2 simultaneously:

$$\Psi_{b,i,g} = B_g(\tanh X_{b,i}), \quad \mathbf{Y} = \underbrace{\Psi_{\text{flat}} \mathbf{C}^\top}_{\text{spline branch}} + \underbrace{\text{silu}(\mathbf{X}) \mathbf{W}^\top}_{\text{base branch}}, \quad \mathbf{C} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}} G}, \quad \mathbf{W} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}, \quad (3)$$

where $\Psi_{\text{flat}} \in \mathbb{R}^{B \times d_{\text{in}} G}$ is Ψ with its last two axes merged. Entry $C_{j,(i,g)}$ is the coefficient c_g of edge $i \rightarrow j$, and $W_{j,i}$ is that edge’s residual weight w_b , so $Y_{b,j} = \sum_i \phi_{i,j}(X_{b,i})$ exactly. Figure 2 traces the tensor shapes, and the forward pass is short enough to show in full:




Figure 1. The KAN-ODE training loop as implemented. The integrator calls the network once per Runge-Kutta stage, so one training epoch on $[0, 3.5]$ evaluates $\mathbf{f}_\theta$ 70s times (35 sampling intervals, two sub-steps each, s stages per step). The loss gradient is obtained by differentiating through every one of those evaluations. Numbered badges mark the three components this report varies: the edge basis (Section 3), the integrator (Sections 3, 5, 7) and the gradient method (Section 4).





Figure 2. Tensor flow through one KDense layer (Equation 3). Only the basis-expansion step differs between the seven basis families studied in Section 3; every basis function has the same signature, $(\mathbf{U}, \mathbf{z}, h) \mapsto \Psi \in \mathbb{R}^{B \times d_{in} \times G}$, so swapping the basis changes one tensor and nothing else.

Parameter count. Each layer holds $d_{out} d_{in} G$ spline coefficients and $d_{out} d_{in}$ residual weights and no bias. For $[2, 10, 2]$ with $G = 5$: layer one ($2 \rightarrow 10$) has $10 \cdot 2 \cdot 5 + 10 \cdot 2 = 100 + 20 = 120$ parameters, layer two ($10 \rightarrow 2$) has $2 \cdot 10 \cdot 5 + 2 \cdot 10 = 100 + 20 = 120$, for the 240 total the base paper reports. The comparison MLPs in

Section 2 are sized to match: the paper’s $[2, 50, 2]$ network has $2 \cdot 50 + 50 + 50 \cdot 2 + 2 = 252$ parameters, and the deeper $[2, 14, 8, 8, 2]$ network used as the working baseline also has $(28 + 14) + (112 + 8) + (64 + 8) + (16 + 2) = 252$.

Why the efficient-kan formulation. The reference implementation, pykan [8], represents every edge with its own spline object and recomputes the grid during training. Inside a Neural ODE this cost is multiplied by the integrator: one epoch of Tsit5 on the training window calls $\mathbf{f}_\theta$ 420 times, and reverse-mode differentiation keeps every one of those calls in the autograd graph. This project’s layer follows the efficient-kan reformulation [1] instead: the knots are a fixed buffer, the basis is evaluated for all edges by one broadcast, and the contraction is one dense matrix product per branch (Equation 3). Each stage therefore adds two GEMMs and a handful of elementwise operations to the graph, independent of the number of edges. The same structure is what makes every basis substituted into Equation 2 (RBF, B-spline, Chebyshev, Lagrange, Newton, IQF, RSWAF) a swappable tensor rather than a swap of object type.

What was reimplemented. The KAN layer, the seven basis functions, the six fixed-step Runge-Kutta integrators, the invariant-enforcing field wrappers used for SIR, and the training loop were written directly for this project rather than imported, so every basis and solver could be instrumented identically. The integrators apply the tableaux of Section 3 verbatim; Tsit5 and DOPRI5 are run at fixed step with their fifth-order weights only, whereas the base paper uses Tsit5’s embedded error estimate for adaptive step control. The base paper’s own implementation compiles its integration loop in Julia’s SciML ecosystem; this project runs the same loop uncompiled in PyTorch, so wall-clock figures reported later are comparable across this project’s own runs but not directly against the paper’s reported times (Section 2.3 quantifies what this costs).

1.3 What the base paper fixed, and what this project varies

The base paper reports two central claims on the Lotka-Volterra predator-prey system: a KAN-ODE with 240 parameters extrapolates Lotka-Volterra’s limit cycle accurately from a short training window, and it reaches a target training loss in roughly 10^4 epochs against a parameter-matched multilayer-perceptron Neural ODE’s 10^5 – a claimed tenfold training-speed advantage at 252 parameters. To obtain these results, the paper commits to three specific design decisions that this project treats as free variables rather than fixed choices: the **integrator** (an adaptive fifth-order Runge-Kutta method, Tsit5), the **basis** (Gaussian RBF), and the **differentiation method** used to compute gradients through the integrator (continuous adjoint sensitivity, motivated by its constant memory cost regardless of trajectory length).

The basis choice is not a detail invented for this project to vary: the original Kolmogorov-Arnold Network paper [8] used the cubic B-spline as its default edge basis, and the base paper switches to Gaussian RBF specifically for the ODE setting, since its derivative never vanishes – relevant once the edge function sits inside a differentiated integrator rather than a static regression. Section 3 restores B-spline and adds five further bases (Chebyshev, Lagrange, Newton, and the IQF and RSWAF kernels), so the trade-off the switch was based on is measured directly rather than assumed.

Each of those three fixed choices maps directly onto one of this report’s studies: Section 4 tests whether the adjoint method’s memory motivation actually holds at this project’s scale; Sections 3, 5, and 7 test the integrator choice, sweeping solver order and, in the stiffness study, solver order jointly with a physical stiffness parameter; and Section 6 tests the basis choice directly, by making the basis itself a learnable, trained quantity rather than a fixed decision made in advance. Section 8 extends the comparison outside the base paper’s own scope entirely, benchmarking KAN-ODE against a sparse-regression alternative built for the same symbolic-discovery goal the RBF basis was partly chosen to support, on both synthetic data and a real, non-mechanistic epidemic curve.

1.4 Why training itself needs to be treated as an object of study

This project uses two testbeds beyond Lotka-Volterra – the SIR epidemic model and a damped pendulum – chosen for simple, well-understood governing equations, so any training failure is attributable to the architecture or optimization rather than to ambiguous target dynamics. The integrator’s interaction with training is not predictable from classical numerical-analysis theory alone, because it sits inside a backpropagation graph rather than a standalone forward solve: a solver that is unconditionally stable classically can still train badly, train to a wrong answer, or fail to generalize. Section 3.5 makes this concrete: a network trained through Forward Euler

learns the *inverse modified equation* of backward error analysis rather than the physical field (its deviation from the true field matches the predicted $\frac{h}{2}\mathbf{J}\mathbf{f}$ correction with cosine similarity 0.92 and magnitude ratio 1.04), and every trained model – KAN or MLP – replaces Lotka-Volterra’s family of neutral orbits with a single attracting limit cycle. Diagnosing behaviour like this means treating training itself as an object of study, not just the final trajectory error.

1.5 Structure of this report

The report follows the order the work was carried out. Section 2 fixes the training recipe, the evaluation protocol and the compute budget, and reproduces the base paper’s KAN-ODE versus MLP-ODE comparison. Section 3 documents the integrators and basis functions from their defining formulas, sweeps six solvers, seven bases, three step sizes and four noise levels, analyses the trained models numerically (order conditions, stability regions, equilibria, invariants, field error), and then diagnoses and fixes two cross-domain training failures. Five further studies build on that fixed baseline: gradient-method cost (Section 4), whether gradient noise explains a solver-order effect (Section 5), a learnable basis blend (Section 6), solver-stiffness interaction (Section 7), and KAN-ODE against sparse regression on synthetic and real data (Section 8). The validation stage (Section 9) retrain five configurations at up to five times the budget and audits the implementation against the paper. Section 10 draws together cross-cutting patterns, and Section 11 closes the report.

2 Baseline Recipe, Evaluation Protocol and Reproduction

The baseline stage reproduces the base paper’s KAN-ODE recipe [6] on the Lotka-Volterra predator-prey system,

$$\dot{u}_1 = \alpha u_1 - \beta u_1 u_2, \quad \dot{u}_2 = \delta u_1 u_2 - \gamma u_2, \quad (\alpha, \beta, \gamma, \delta) = (1.5, 1.0, 3.0, 1.0), \quad (4)$$

using a two-layer Kolmogorov-Arnold Network, $[2, 10, 2]$, with $G = 5$ Gaussian radial-basis-function grid points per edge (240 parameters), trained on $t \in [0, 3.5]$ (36 points) and scored strictly on $t \in (3.5, 14.0]$ (105 unseen points). Extrapolation is scored this way, on points the training loss never touched, because it is the only meaningful test of whether the network learned the underlying dynamics rather than memorized the shape of the training window; a model that only reproduces the 36 training points could still be badly wrong about the system it is meant to describe.

Two properties of Equation 4 are used repeatedly below as exact, model-independent checks. It has a single interior equilibrium $\mathbf{u}^* = (\gamma/\delta, \alpha/\beta) = (3, 1.5)$, where the Jacobian has purely imaginary eigenvalues $\pm i\sqrt{\alpha\gamma} = \pm 2.121i$. And it conserves the first integral

$$H(u_1, u_2) = \delta u_1 - \gamma \ln u_1 + \beta u_2 - \alpha \ln u_2, \quad \frac{dH}{dt} = 0, \quad (5)$$

so every orbit is a closed level curve of H and the phase plane is filled with a continuous family of neutral cycles. The orbit through $\mathbf{u}(0) = (1, 1)$ has $H_0 = 2$ and period $T = 3.3175$ (measured from a DOP853 solve at tolerance 10^{-12}). The training window therefore covers 1.06 periods, the scored window 3.2 further periods, and the full horizon $[0, 14]$ about 4.2.

Two additional test systems were carried forward alongside Lotka-Volterra, chosen because both have closed-form governing equations that are textbook-simple and well understood, so that any training failure can be attributed to the architecture or the optimization process rather than to an ambiguous target dynamics.

The **SIR compartmental epidemic model** partitions a fixed population into susceptible (S), infected (I), and recovered (R) compartments,

$$\dot{S} = -\beta SI, \quad \dot{I} = \beta SI - \gamma I, \quad \dot{R} = \gamma I, \quad (\beta, \gamma) = (0.35, 0.1), \quad (6)$$

with a built-in algebraic invariant: the three compartments sum to a constant, $S + I + R = 1$, because every term in the field cancels when summed. This invariant is exact and independent of the parameter values, so any model that violates it has learned a vector field whose structure is wrong, not merely inaccurate.

The **damped pendulum** evolves angular displacement θ and angular velocity ω under gravity and viscous friction,

$$\dot{\theta} = \omega, \quad \dot{\omega} = -\frac{g}{L} \sin \theta - \mu \omega, \quad (g/L, \mu) = (9.81, 0.5), \quad (7)$$

where $\mu > 0$ guarantees strict energy dissipation: the mechanical energy $E = \frac{1}{2}mL^2\omega^2 + mgL(1 - \cos \theta)$ satisfies $dE/dt = -mL^2\mu\omega^2 \leq 0$. A model can fit the trajectory accurately and still violate this monotonic decay, making it a genuinely independent physical check.

2.1 Training recipe

Table 1 lists every setting shared by the Lotka-Volterra runs; each experiment below changes exactly one of them. Ground truth is generated with SciPy’s DOP853 at relative and absolute tolerance 10^{-12} , so the reference trajectory is accurate to far below any error reported here.

Table 1. Shared training recipe for all Lotka-Volterra runs.

Setting	Value	Note
Network	KAN [2, 10, 2], $G = 5$, 240 parameters	knots uniform on $[-1, 1]$, $h = 0.5$
Edge function	Equation 2, tanh normaliser, SiLU residual	Glorot-uniform initialisation of C and W
Integrator	Tsit5, fixed step	2 sub-steps per sample: $h = \Delta t/2 = 0.05$
Sampling interval	$\Delta t = 0.1$	36 training points, 141 on $[0, 14]$
Loss	mean squared trajectory error on $[0, 3.5]$	no regularisation ($\lambda_{\text{act}} = \lambda_{\text{ent}} = 0$)
Optimiser	Adam, learning rate 2×10^{-3}	constant learning rate, no schedule
Gradient clipping	global norm ≤ 1.0	non-finite gradients skip the step (Section 3)
Budget	10,000 epochs, full batch	50,000 in the validation stage (Section 9)
Model selection	checkpoint with minimum training MSE	the extrapolation window is never used for selection
Precision, seed	float32, seed 42	one seed per configuration ($N = 1$)

2.2 Evaluation protocol

This stage established the evaluation protocol used throughout the remainder of this report (Figure 3). Model selection is based on minimum *training*-window loss only, and every accuracy number is computed strictly outside the training horizon, so that a number cannot improve simply by widening what counts as fit.




Figure 3. Time axis of the Lotka-Volterra protocol. Tick marks above the bar are multiples of the true period. The far-horizon window $(14, 28]$ is used only for the stress tests of Sections 3 and 9; no model is retrained or reselected for it.

Writing $\hat{\mathbf{u}}_k$ for the model’s prediction and $\mathbf{u}_k$ for the truth at the $N_e = 105$ extrapolation samples, the reported metrics are

$$\text{MSE} = \frac{1}{2N_e} \sum_k \|\hat{\mathbf{u}}_k - \mathbf{u}_k\|_2^2, \quad R^2 = 1 - \frac{\sum_k \|\hat{\mathbf{u}}_k - \mathbf{u}_k\|_2^2}{\sum_k \|\mathbf{u}_k - \bar{\mathbf{u}}\|_2^2}, \quad \text{rel. } L_2 = \frac{(\sum_k \|\hat{\mathbf{u}}_k - \mathbf{u}_k\|_2^2)^{1/2}}{(\sum_k \|\mathbf{u}_k\|_2^2)^{1/2}}, \quad (8)$$

where $\bar{\mathbf{u}}$ is the componentwise mean of the true extrapolation trajectory. $R^2 = 1$ is a perfect match, $R^2 = 0$ is no better than predicting that mean, and a negative value is worse than that trivial baseline. Two further quantities describe the trained field rather than its trajectory. The **Lipschitz estimate** is

$$\hat{L} = \max_{k \leq 36} \frac{\|\mathbf{f}_\theta(\mathbf{u}_k + \varepsilon \mathbf{v}_k) - \mathbf{f}_\theta(\mathbf{u}_k)\|_2}{\varepsilon}, \quad \varepsilon = 10^{-3}, \quad \|\mathbf{v}_k\|_2 = 1 \text{ random}, \quad (9)$$

a finite-difference sample of the local slope of $\mathbf{f}_\theta$ at the training states, so it is a *lower* bound on the field’s true Lipschitz constant and a measure of how sharply the learned field bends. The **cost** of a solver is its number of

function evaluations: an s -stage method on the training window costs $\text{NFE} = 35 \times 2 \times s = 70s$ calls of $\mathbf{f}_\theta$ per epoch, which, because reverse-mode differentiation unrolls every call, is also what training wall-clock scales with.

2.3 Computational budget: why 10,000 epochs

The base paper trains for 10^5 epochs. Table 2 records what that would cost here. The paper’s integration loop is compiled by Julia’s SciML stack; this project’s is interpreted PyTorch on CPU, where every Runge-Kutta stage is a separate Python-level call with its own autograd bookkeeping. On the development machine one Tsit5/RBF epoch takes 0.17 s and one Tsit5/B-spline epoch 0.67 s; the cloud CPUs used for the long runs of Section 9 were measured at $2.2\text{--}2.5\times$ slower for every KAN configuration.

Table 2. Measured per-epoch cost, local development CPU against cloud (Kaggle) CPU, each averaged over the full run, and the projected cost of the paper’s 10^5 -epoch budget.

Configuration	Local s/epoch	Cloud s/epoch	Ratio	10^5 epochs, local	10^5 epochs, cloud
Euler / RBF	0.0254	0.0627	$2.47\times$	0.7 h	1.7 h
Tsit5 / RBF (reference)	0.1705	0.4033	$2.37\times$	4.7 h	11.2 h
Tsit5 / B-spline	0.6680	1.5438	$2.31\times$	18.6 h	42.9 h
MLP-ODE, SiLU [2, 14, 8, 8, 2]	0.2856	0.2722	$0.95\times$	7.9 h	7.6 h
MLP-ODE, tanh [2, 50, 2]	0.0736	0.1595	$2.17\times$	2.0 h	4.4 h
All 24 Lotka-Volterra configurations	14.2 CPU-h at 10^4 epochs			142 h	≈ 336 h

A 10^5 -epoch version of the 24-configuration Lotka-Volterra sweep would need roughly 142 CPU-hours locally or about two weeks of cloud CPU-time, before the pendulum, SIR and follow-on studies. The sweep was therefore run at 10^4 epochs, a budget at which every noise-free solver and basis configuration except Euler and four bases (Lagrange, IQF, RSWAF, Newton) reaches 10^{-4} training MSE (Table 11). That choice is only safe if the conclusions drawn at 10^4 epochs survive a longer budget, which is exactly what Section 9 tests.

2.4 Baseline reproduction: KAN-ODE against parameter-matched MLP-ODEs

Two MLP vector fields with 252 parameters were trained under the identical recipe: the paper’s literal baseline, one hidden layer of 50 tanh units, and a deeper [2, 14, 8, 8, 2] SiLU network that trains reliably under this protocol. Table 3 and Figures 4 and 5 compare them with the KAN-ODE. A fourth column, added after auditing the paper’s own code (Section 9.3), retrains the identical literal baseline under the paper’s own learning rate and initialization scale instead of this project’s default; Figures 4 and 5 plot that run in place of the default-setting tanh MLP.

The fourth column shows that the paper’s literal architecture was never the problem: given the paper’s own learning rate and initialization scale, the identical tanh network reaches $R^2 = 0.997$, a discussion deferred to Section 9 where its 5×10^4 -epoch continuation is also reported. The next two paragraphs compare KAN-ODE against the working SiLU baseline, columns one and two.

Two findings, one confirming the paper and one contradicting it. **KAN-ODE extrapolates $5.4\times$ better** from a training loss only 7% lower, reproducing the paper’s accuracy claim. The checkpoints show where that gap comes from. Both fields are equally accurate on the orbit (2.84% against 2.82% mean field error), and both hold the invariant to about 3%, but the MLP’s orbit period is off by 8.0×10^{-4} against the KAN’s 1.7×10^{-5} , a $48\times$ larger period error. On a periodic orbit a period error ε_T becomes a phase lag that grows linearly, $\approx 2\pi\varepsilon_T t/T$, so the trajectory error of the MLP keeps accumulating while the KAN’s stays near its training-window level. The extrapolation gap is a phase-accuracy gap rather than a field-accuracy gap. **The MLP converges faster, not slower:** it reaches 10^{-3} training MSE at epoch 1297 against the KAN’s 5282, $4.1\times$ sooner. This contradicts the paper’s headline efficiency claim on this same system (a 10^4 -epoch KAN-ODE against a 10^5 -epoch MLP-ODE). Under this project’s protocol, at matched parameter count and optimizer settings, the MLP was the faster model to train; KAN-ODE won on extrapolation, the paper’s other headline claim.




Figure 4. Training MSE per epoch for KAN-ODE and the paper’s parameter-matched tanh MLP-ODE, trained with the paper’s own learning rate and initialization (fourth column of Table 3). KAN-ODE descends steadily and reaches 10^{-4} at epoch 9617; the tanh MLP reaches 10^{-2} at epoch 3461 but never 10^{-3} within 10^4 epochs.




Figure 5. Predicted against true populations out to $t = 28$ for the three models, with no retraining. KAN-ODE and SiLU-MLP are visually indistinguishable from the truth. The paper’s tanh MLP under its own learning rate and initialization (bottom, fourth column of Table 3) follows the training window and the first periods closely, then runs slightly slow and lags visibly by the far window $[14, 28]$.

Table 3. KAN-ODE against parameter-matched MLP-ODEs at 10^4 epochs. Epoch counts are the first epoch at which the training loss falls below the threshold. The last three rows are computed from the saved checkpoints: field error is $\|\mathbf{f}_\theta - \mathbf{f}\|_2$ averaged over the true orbit on $(3.5, 14]$ and divided by the mean $\|\mathbf{f}\|_2 = 4.82$ there; the period is measured from zero-crossings of the predicted prey signal over $[0, 28]$; H drift is $\max_t |H(\hat{\mathbf{u}}(t)) - H_0|/H_0$ on $[0, 28]$. The fourth column is the identical literal baseline of the third, retrained under the paper’s own learning rate and initialization scale (Section 9.3); its 5×10^4 -epoch continuation is in Table 23.

Metric	KAN-ODE (RBF)	MLP-ODE (SiLU)	MLP-ODE (tanh, default lr/init)	MLP-ODE (tanh, paper’s lr/init)
Parameters	240	252	252	252
Training MSE	8.83×10^{-5}	9.50×10^{-5}	1.08×10^0	1.11×10^{-3}
Extrapolation MSE	8.92×10^{-5}	4.77×10^{-4}	1.69×10^0	9.65×10^{-3}
Extrapolation R^2	0.99997	0.99984	0.4215	0.99669
Relative L_2 error	0.33%	0.76%	45.2%	3.42%
Lipschitz estimate $\hat{L}$	6.91	5.07	38.22	9.23
Epochs to 10^{-2} / 10^{-3} / 10^{-4}	3112 / 5282 / 9617	741 / 1297 / 9145	never	3461 / never / never
Wall-clock, 10^4 epochs	1705 s	2856 s	736 s	1052 s
Field error on the orbit, (3.5, 14]	2.84%	2.82%	–	7.03%
Relative period error	1.7×10^{-5}	8.0×10^{-4}	–	3.1×10^{-3}
Max. first-integral drift	3.1%	2.8%	–	22.7%

Why the paper’s literal MLP fails to train under this project’s default optimizer settings. The $[2, 50, 2]$ tanh network never gets below a training MSE of 1.08 at this project’s default learning rate (2×10^{-3}) and initialization (Glorot, gain 1.0). A controlled 2000-epoch factorial run, holding those optimizer settings fixed, separates the two things that differ between it and the working baseline (Table 4).

Table 4. Controlled breakdown of why the paper’s literal MLP baseline architecture, $[2, 50, 2]$ with hyperbolic tangent (tanh) activation, fails to train at this project’s default learning rate and initialization scale. All four runs use 2000 epochs and Glorot initialization at gain 1.0; none varies the initialization scale (see the paragraph following this table).

Run	Architecture	Activation	Learning rate	Training MSE	Extrapolation MSE	R^2
A	$[2, 50, 2]$	tanh	2×10^{-3}	1.21×10^0	1.98×10^0	0.319
B	$[2, 50, 2]$	tanh	5×10^{-4}	1.85×10^0	9.23×10^0	–2.169
C	$[2, 14, 8, 8, 2]$	SiLU	2×10^{-3}	4.28×10^{-4}	6.97×10^{-4}	0.99976
D	$[2, 50, 2]$	SiLU	2×10^{-3}	3.62×10^{-3}	2.68×10^2	–91.06

Runs A and D share the architecture and differ only in activation: the training MSE falls from 1.21 to 3.62×10^{-3} , a $335\times$ improvement, so at these optimizer settings, activation decides whether this architecture fits the data at all. Runs C and D share the SiLU activation and differ only in depth: the shallow network fits the window but its extrapolation error explodes to 2.68×10^2 , against 6.97×10^{-4} for the deep one, so depth decides whether a model that trains also extrapolates. Lowering the learning rate alone (run B) makes the tanh network worse, not better, and five times the epoch budget under these same settings does not rescue it either: at 5×10^4 epochs, training MSE moves only from 1.08 to 1.01 and extrapolation gets *worse* (Table 23).

The actual cause: learning rate and initialization scale, not architecture. None of runs A–D varies the initialization scale, and the base paper’s own published code turned out to use a combination this factorial never tested: a learning rate of 10^{-2} (five times this project’s default) together with weights initialized near zero (Glorot divided by 10^5), for the identical $[2, 50, 2]$ tanh architecture. Retrained with only those two settings changed, the literal architecture trains (Table 3, fourth column): at 10^4 epochs it reaches a training MSE of 1.11×10^{-3} (against 1.08 before) and an extrapolation MSE of 9.65×10^{-3} , $R^2 = 0.997$ (against $R^2 = 0.42$ before), with a Lipschitz estimate of 9.23 against 38.2; at 5×10^4 epochs it reaches 5.43×10^{-5} training MSE and $R^2 = 0.9998$ on extrapolation, with zero non-finite gradient steps and no clipping in either run. Section 9 reports the full

comparison against KAN-ODE and the SiLU MLP at both budgets, and updates the implementation audit accordingly. The paper’s literal recipe was never architecturally incompatible with this system: this project’s default optimizer settings, tuned around the KAN-ODE side of the comparison, were the wrong settings for a wide, near-linear-at-initialization tanh network, and Table 4’s factorial measured sensitivity to activation *at those settings* rather than an architecture-level limit.

3 Benchmark Ablation and Cross-Domain Stability Analysis

Before extending KAN-ODE to new systems, two questions needed settling: which solver and basis choices this architecture actually prefers, and whether the baseline recipe (Section 2) transfers unchanged to systems beyond Lotka-Volterra. Both were answered empirically, and the second uncovered two genuine training failures that required a mechanistic diagnosis before they could be fixed. This section first documents the integrators and bases from their defining formulas and verifies the implementations numerically, then runs the ablations, then analyses what the trained vector fields actually learned, and finally turns to the two cross-domain failures.

3.1 The six integrators, and a check that each has its design order

Every integrator in Table 7 is an explicit Runge-Kutta method: one step of size h evaluates s stage slopes and combines them,

$$\mathbf{k}_i = \mathbf{f}_\theta\left(\mathbf{u}_n + h \sum_{j=1}^{i-1} a_{ij} \mathbf{k}_j\right), \quad i = 1, \dots, s, \quad \mathbf{u}_{n+1} = \mathbf{u}_n + h \sum_{i=1}^s b_i \mathbf{k}_i, \quad (10)$$

so a method is fully specified by its Butcher tableau $(A, \mathbf{b}, \mathbf{c})$ with $c_i = \sum_j a_{ij}$. All six were implemented directly for this project rather than called from an external ODE library. Writing $\mathbf{u}_n$ for the state at step n , the four low-order methods are

$$\text{Forward Euler (order 1): } \mathbf{u}_{n+1} = \mathbf{u}_n + h \mathbf{f}_\theta(\mathbf{u}_n) \quad (11)$$

$$\text{Midpoint (order 2): } \mathbf{k}_1 = \mathbf{f}_\theta(\mathbf{u}_n), \quad \mathbf{u}_{n+1} = \mathbf{u}_n + h \mathbf{f}_\theta\left(\mathbf{u}_n + \frac{h}{2} \mathbf{k}_1\right) \quad (12)$$

$$\text{Heun’s method (order 2): } \mathbf{k}_1 = \mathbf{f}_\theta(\mathbf{u}_n), \quad \mathbf{k}_2 = \mathbf{f}_\theta(\mathbf{u}_n + h \mathbf{k}_1), \quad \mathbf{u}_{n+1} = \mathbf{u}_n + \frac{h}{2}(\mathbf{k}_1 + \mathbf{k}_2) \quad (13)$$

$$\text{Classical RK4 (order 4): } \mathbf{u}_{n+1} = \mathbf{u}_n + \frac{h}{6}(\mathbf{k}_1 + 2\mathbf{k}_2 + 2\mathbf{k}_3 + \mathbf{k}_4) \quad (14)$$

with the classical RK4 stages $\mathbf{k}_1 = \mathbf{f}_\theta(\mathbf{u}_n)$, $\mathbf{k}_2 = \mathbf{f}_\theta(\mathbf{u}_n + \frac{h}{2} \mathbf{k}_1)$, $\mathbf{k}_3 = \mathbf{f}_\theta(\mathbf{u}_n + \frac{h}{2} \mathbf{k}_2)$, $\mathbf{k}_4 = \mathbf{f}_\theta(\mathbf{u}_n + h \mathbf{k}_3)$. DOPRI5 and Tsit5 are both seven-stage 5(4) embedded pairs whose seventh stage exists only to supply the fourth-order error estimate (the “first same as last” stage); its fifth-order weight is $b_7 = 0$. Run at fixed step, as here, the seventh evaluation contributes nothing to the update, so the implementation skips it and both methods cost *six* evaluations per step. Their coefficients are chosen so that all 17 order conditions up to order five hold; Tsitouras [9] additionally minimised the size of the sixth-order error terms. Table 5 lists the Tsit5 coefficients exactly as implemented, and Figure 6 shows where each method samples the field within one step. Every extra stage costs one additional evaluation of $\mathbf{f}_\theta$, which is why function evaluations per epoch, not solver order alone, is the operative cost measure.

Table 5. Fixed-step Tsit5 tableau as implemented (first six stages, rounded to six significant figures; the code carries 16). Row i lists c_i and $a_{i1}, \dots, a_{i,i-1}$; the last row lists the solution weights b_i .

0						
0.161	0.161					
0.327	−0.00848066	0.335481				
0.9	2.89715	−6.35945	4.36230			
0.980026	5.32586	−11.7489	7.49554	−0.0924951		
1	5.86146	−12.9210	8.15937	−0.0715850	−0.0282691	
$\mathbf{b}^\top$	0.0964608	0.01	0.479890	1.37901	−3.29007	2.32471

Verifying the implementations. Two independent checks confirm that each integrator is the method it claims to be. First, the order conditions: a Runge-Kutta method has order p exactly when $\mathbf{b}^\top \Phi(\tau) = 1/\gamma(\tau)$ for every rooted tree τ with at most p vertices, where $\Phi(\tau)$ is the tree’s elementary weight built from A and $\gamma(\tau)$ its




Figure 6. Stage nodes $t_n + c_i h$ and solution weights b_i of the six implemented methods, drawn from the coefficients in the code. Heun and Midpoint use the same number of evaluations but sample at opposite places: Heun averages the two ends of the step, Midpoint uses only a slope at its centre. Tsit5 reaches order five with large weights of both signs ($b_5 = -3.29$), a cancellation that is harmless in exact arithmetic and is what buys its small error constant.

density [5]. Enumerating all 37 rooted trees up to six vertices and evaluating each condition directly on the implemented coefficients gives the verified order and the principal error norm $A^{(p+1)} = (\sum_{|\tau|=p+1} [(\mathbf{b}^\top \Phi(\tau) - 1/\gamma(\tau))/\sigma(\tau)]^2)^{1/2}$, the size of the leading local-error term. Second, a work-precision study on the *true* Lotka-Volterra field: integrating from $(1, 1)$ to $t = 3.5$ at eight step sizes and comparing against the DOP853 reference gives the observed order $\hat{p}$ as the slope of $\log \|\text{error}\|$ against $\log h$ (Figure 7). Table 6 collects both.

Table 6. Order verification of the implemented integrators. “Max residual” is the largest $|\mathbf{b}^\top \Phi(\tau) - 1/\gamma(\tau)|$ over all trees up to the design order ($1 + 1 + 2 + 4 + 9 = 17$ conditions at order five); $A^{(p+1)}$ is the principal error norm; $\hat{p}$ is the slope fitted to the work-precision data for $h \leq 0.1$ above the round-off floor.

Method	Stages s	Verified order p	Conditions checked	Max residual through order p	$A^{(p+1)}$	Observed $\hat{p}$
Forward Euler	1	1	1	0	5.00×10^{-1}	1.00
Heun	2	2	2	0	1.86×10^{-1}	2.10
Midpoint	2	2	2	0	1.72×10^{-1}	2.07
Classical RK4	4	4	8	1.1×10^{-16}	1.45×10^{-2}	3.93
DOPRI5 (6 of 7)	6	5	17	2.8×10^{-16}	3.99×10^{-4}	5.66
Tsit5 (6 of 7)	6	5	17	1.4×10^{-14}	1.39×10^{-4}	5.95

Every method satisfies its order conditions to round-off and shows its design order in practice; the fifth-order pair shows slightly super-nominal slopes over this range because its error is still pre-asymptotic before it hits round-off. The principal error norms confirm Tsitouras’s design goal from the code’s own coefficients: Tsit5’s leading error term is $2.9\times$ smaller than DOPRI5’s. The table also sets up a point used below: Heun and Midpoint differ only in their third-order error terms, and on those Midpoint’s is marginally *smaller* (0.172 against 0.186).

Linear stability. Applied to $\dot{u} = \lambda u$, every explicit Runge-Kutta method multiplies the state by the stability polynomial $R(z) = 1 + z \mathbf{b}^\top (I - zA)^{-1} \mathbf{1}$, $z = h\lambda$, per step. Heun and Midpoint share $R(z) = 1 + z + z^2/2$, so




Figure 7. Global error at $t = 3.5$ of each implemented integrator applied to the true Lotka-Volterra field. (a) Against step size, with fitted slopes; the dashed line is the training step $h = 0.05$. (b) Against the number of field evaluations, the quantity that sets training cost. At $h = 0.4$ both second-order methods blow up, and the fifth-order methods reach the round-off floor near 10^{-13} .

they are indistinguishable on any linear problem. Figure 8a draws the boundaries $|R(z)| = 1$ and overlays $h\lambda$ for the eigenvalues of the true Jacobian and of the trained Tsit5 KAN's Jacobian $J_\theta = \partial \mathbf{f}_\theta / \partial \mathbf{u}$, evaluated at all 141 states of the orbit. The learned eigenvalues track the true ones closely ($\max |\lambda| = 3.61$ against 3.83), and at $h = 0.05$ every $h\lambda$ lies within 0.19 of the origin, far inside every region. Stability in the classical sense is therefore never at issue at the training step; what differs between methods is how accurately R reproduces e^z near the origin. For a purely oscillatory mode $z = i\omega h$, Figure 8b plots the per-step amplitude error $|R(i\omega h)| - 1$: at Lotka-Volterra's linear frequency $\omega_0 = \sqrt{\alpha\gamma}$, a linear oscillation integrated over one period gains 39% amplitude under Euler and 0.094% under either second-order method, and loses 5.8×10^{-7} , 2.3×10^{-8} and 3.5×10^{-9} under RK4, DOPRI5 and Tsit5.




Figure 8. (a) Stability boundaries $|R(z)| = 1$ computed from each implemented tableau, with $h\lambda$ for the true Jacobian (dots) and the learned Jacobian of the Tsit5-trained KAN (crosses) at every state on the orbit, $h = 0.05$; the inset enlarges the neighbourhood of the origin. (b) Per-step amplitude error for an undamped oscillation; the dotted line marks $h\omega_0$ at the training step.

The linear picture is borne out on the true nonlinear field. Integrating the exact Lotka-Volterra field at the training discretisation and scoring it with the same extrapolation metric as the trained models isolates each solver's own truncation error (last column of Table 7): it is 17.6 for Euler, whose trajectory leaves the positive quadrant at $t = 19.4$, $1.3\text{--}3.0 \times 10^{-4}$ for the second-order pair, and below 5×10^{-9} for the rest. The first integral of Equation 5 tells the same story (Figure 18a): by $t = 28$ it has drifted by 7.5×10^{-4} and 4.5×10^{-4} (relative) under Heun and Midpoint, 3.0×10^{-5} under RK4, and about 10^{-7} under the fifth-order pair.

3.2 Solver ablation

Table 7 sweeps the six integrators at fixed RBF basis; Table 8 below sweeps seven basis representations at fixed Tsit5 solver. Both use the protocol of Section 2.2. The Lipschitz estimate of the trained field (Equation 9) is reported alongside accuracy because it measures how smooth or jagged the basis expansion in Equation 2 ended up, independent of the trajectory error.

Table 7. Solver ablation (fixed RBF basis, $h = 0.05$, 240 parameters, 10^4 epochs). The last column integrates the *true* field with the same solver and step and scores it identically: the solver’s own truncation error, with no learning involved.

Solver	Order	NFE / epoch	Training MSE	Extrap. MSE	Extrap. R^2	$\hat{L}$	Wall- clock	True-field truncation MSE
Forward Euler	1	70	1.38×10^{-4}	2.51×10^{-3}	0.99914	8.55	254 s	1.8×10^1
Heun RK2	2	140	8.32×10^{-5}	3.10×10^{-4}	0.99989	7.07	499 s	1.3×10^{-4}
Midpoint	2	140	8.86×10^{-5}	8.95×10^{-5}	0.99997	6.96	498 s	3.0×10^{-4}
Classical RK4	4	280	9.41×10^{-5}	9.29×10^{-5}	0.99997	6.88	1019 s	4.7×10^{-9}
DOPRI5	5	420	9.28×10^{-5}	9.21×10^{-5}	0.99997	6.89	1699 s	6.7×10^{-14}
Tsit5	5	420	8.83×10^{-5}	8.92×10^{-5}	0.99997	6.91	1712 s	4.7×10^{-14}

Truncation error dominates only at order $p = 1$: Euler’s extrapolation MSE is $28\times$ Tsit5’s, and it is the only solver that fails to reach 10^{-4} training loss within budget. Its trained field is also the least smooth of the six ($\hat{L} = 8.55$, the highest in the table), consistent with a systematic bias forcing the network to fit a sharper correction than the smoother integrators require. Above order 2 the ranking collapses into a 3.8% spread with no monotonic order dependence – at $h = 0.05$, integrator truncation error has fallen below the optimization noise floor. Midpoint matches Tsit5’s accuracy within 0.3% at one third of the function evaluations and 29% of the wall-clock, making it the cost-normalized optimum on this problem.

The last column is the striking one. Integrated with the *true* field, Euler’s own error on the extrapolation window is 17.6; the Euler-trained network, integrated with Euler, scores 2.5×10^{-3} , nearly four orders of magnitude better than the exact field under the same solver. The same holds, more mildly, for Midpoint, whose trained model (8.95×10^{-5}) beats the exact field under Midpoint (3.0×10^{-4}). The network is not learning $\mathbf{f}$; it is learning whatever field makes *this solver’s* output match the data. Two experiments make that precise.

The Euler model learns the inverse modified equation. Backward error analysis states that Forward Euler applied to a field $\mathbf{g}$ follows, to second order, the modified field $\mathbf{g} - \frac{h}{2}J_{\mathbf{g}}\mathbf{g} + O(h^2)$. For Euler’s output to reproduce the exact flow, the network must therefore learn

$$\mathbf{g}_h = \mathbf{f} + \frac{h}{2}J_{\mathbf{f}}\mathbf{f} + O(h^2). \quad (15)$$

On the 36 training states, the Euler-trained field differs from the true field by 9.6% (relative L_2), almost exactly the size of the predicted correction $\frac{h}{2}J_{\mathbf{f}}\mathbf{f}$ (9.2%). The learned correction $\mathbf{f}_{\theta} - \mathbf{f}$ and the predicted one point the same way (cosine similarity 0.92) with a magnitude ratio of 1.04; measured against $\mathbf{g}_h$ instead of $\mathbf{f}$, the Euler model’s error falls to 3.8%, comparable to the 3.0% field error of the Tsit5-trained model. The Euler model is an accurate model of Equation 15.

Cross-solver transfer. If each model has absorbed its solver’s error, it should degrade when integrated with a different solver. Figure 9 evaluates every trained model under every integrator. The Euler model is $70\times$ worse under any other solver (0.17–0.18), and every other model is $1000\text{--}7200\times$ worse under Euler. The second-order models lose $9\text{--}17\times$ when integrated with a fourth- or fifth-order method, and the fourth- and fifth-order models lose $14\text{--}17\times$ under a second-order one. Within each family the models are interchangeable: RK4, DOPRI5 and Tsit5 models agree to three significant figures under all three high-order solvers, so these three have learned fields that coincide up to their $O(h^4)$ differences. The practical rule is that a KAN-ODE trained with a low-order solver must be deployed with that same solver, and that a trained field is a statement about the physics only once the training solver’s truncation error is below the model error.




Figure 9. Extrapolation MSE of each trained model (rows) integrated with each solver (columns), $h = 0.05$. The diagonal (bold) is the configuration it was trained in.




Figure 10. Training wall-clock against extrapolation MSE for the solver and basis sweeps. The dashed step is the Pareto front: Euler, Midpoint, Tsit5 and B-spline. Midpoint dominates RK4 and DOPRI5 outright.

Midpoint against Heun. Midpoint extrapolates $3.5\times$ better than Heun at identical cost and order. A tempting explanation is that Heun’s end-of-step sampling adds numerical damping. The numbers above rule that out: the two methods have the same stability polynomial, so they damp identically on linear modes; Heun is actually the *more* accurate integrator of the true field (1.3×10^{-4} against 3.0×10^{-4} truncation MSE); and the third-order error norms differ by only 8%. The gap belongs to the trained models, not to the integrators. It shows up as phase accuracy: the Heun model’s orbit period is off by 4.7×10^{-4} (relative), the Midpoint model’s by 3.5×10^{-6} (Table 10), and its error envelope grows ten times faster over $(3.5, 28]$ (5.9×10^{-3} against 5.7×10^{-4} per unit time). With a single seed per configuration this ordering could still be an initialisation effect; the solver-level explanation is refuted either way.




Figure 11. Phase portraits of the Euler-, Midpoint- and Tsit5-trained RBF models on $[0, 14]$ (all three plots are titled “KAN-ODE (RBF)” because the basis is shared). Each model is integrated with its own training solver. All three stay on the true orbit to plotting accuracy, including Euler, whose $28\times$ larger extrapolation MSE is phase error along the orbit rather than a visible drift off it.

3.3 Basis functions: definitions, conditioning and locality

Writing z_g for the g -th of G evenly spaced grid centers with spacing $h = 2/(G - 1)$, and $u = \tanh(x) \in [-1, 1]$ the normalized edge input, the seven bases compared here are, exactly as implemented,

$$B_g^{\text{RBF}}(u) = \exp\left(-\frac{1}{2}\left(\frac{u-z_g}{h}\right)^2\right), \quad B_g^{\text{IQF}}(u) = \frac{1}{1 + \left(\frac{u-z_g}{h}\right)^2}, \quad B_g^{\text{RSWAF}}(u) = 1 - \tanh^2\left(\frac{u-z_g}{h}\right), \quad (16)$$

$$B_g^{\text{Cheb}}(u) = T_{g-1}(u) = \cos((g-1)\arccos u), \quad B_g^{\text{Lag}}(u) = \prod_{j \neq g} \frac{u - z_j}{z_g - z_j}, \quad B_g^{\text{Newt}}(u) = \prod_{j < g} (u - z_j), \quad (17)$$

plus the cubic B-spline B_g^{spline} , built by the Cox–de Boor recursion from degree-0 indicator functions by repeated linear blending, $B_{g,0}(u) = \mathbf{1}[t_g \leq u < t_{g+1}]$ and $B_{g,p}(u) = \frac{u-t_g}{t_{g+p}-t_g}B_{g,p-1}(u) + \frac{t_{g+p+1}-u}{t_{g+p+1}-t_{g+1}}B_{g+1,p-1}(u)$, on the clamped knot vector $\mathbf{t} = (-1, -1, -1, -1, 0, 1, 1, 1, 1)$ for $G = 5, k = 3$. The clamping makes the five splines a partition of unity with $B_1(-1) = B_5(1) = 1$, and it is what gives the spline compact support: $B_{g,k}(u) = 0$ outside the $k + 2$ knots the recursion references. Gaussian RBF, IQF and RSWAF are local but never exactly zero; Chebyshev, Lagrange and Newton have no locality at all – every coefficient c_g influences $\phi(x)$ at every x . The residual term $w_b \text{silu}(x)$ in Equation 2 is the same for every basis choice, since it is a fixed part of the edge parameterization, not one of the substituted terms above.

Figure 12 draws the seven families from the code itself and reports the 2-norm condition number $\kappa_2(\Psi)$ of the $400 \times G$ design matrix $\Psi_{ng} = B_g(u_n)$ on a uniform sample of $[-1, 1]$, the quantity that governs how well-posed fitting the coefficients c_g is. At $G = 5$ Chebyshev (2.5), Lagrange (3.7) and B-spline (5.5) are well conditioned, the three kernels sit at 8–34, and Newton is worst at 54. The growth with G separates them further: at $G = 12$ Chebyshev and B-spline are still below 6, RBF and Lagrange reach 69 and 63, and Newton reaches 2.4×10^4 , growing exponentially – the nested products $\prod_{j < g} (u - z_j)$ become nearly collinear columns.




Figure 12. The seven implemented basis families on the $G = 5$ grid (vertical lines are the knots z_g), with the $g = 4$ function highlighted and the 2-norm condition number of the design matrix in each title. The blue band is the interval $u \in [0.25, 1]$ that first-layer inputs actually reach on the Lotka-Volterra orbit. The last panel shows κ_2 against grid size; dashed and dotted grey lines are IQF and RSWAF.

How much of the grid the data uses. Lotka-Volterra’s states are positive, between 0.26 and 6.9 on the orbit, so after the tanh normaliser the first layer sees only $u \in [0.25, 1.0]$ – the right 37.5% of its grid. Measured on the trained RBF model, the two leftmost basis functions ($z_1 = -1$, $z_2 = -0.5$) carry 0.4% and 3.6% of the first layer’s total basis activation; for the B-spline the leftmost carries exactly zero and the next 1.7%. The second layer, whose inputs are the signed hidden activations, uses its grid evenly (every basis function 18–26%). Roughly a third of the first layer’s 100 spline coefficients are therefore nearly inert on this system. Figure 13 shows what locality means against that data interval: raising the coefficient of z_2 changes the spline edge by at most 0.021δ on the data, the RBF edge by 0.32δ , and the Chebyshev edge by the full δ .

B-spline edges out RBF on accuracy after a boundary-knot fix, at $3.9\times$ RBF’s wall-clock time and a visibly higher Lipschitz estimate (8.44 vs. 6.91) – its compact local support lets it fit sharper local features than RBF’s smooth, slowly decaying support without being penalized elsewhere in the domain. The two bases are statistically tied on error (6% apart, inside single-seed noise), which makes RBF the better accuracy-per-second choice. Polynomial global-support bases (Chebyshev, Lagrange) fit the training window well and extrapolate 4–5 $\times$ worse – the expected signature of a basis whose every coefficient influences the whole domain rather than a local region of it. Conditioning alone does not predict the ranking: Chebyshev is the best-conditioned basis




Figure 13. Locality of one coefficient update, $c_2 \leftarrow c_2 + \delta$, for three basis families (uniform-knot cubic spline shown for clarity). Since ϕ is linear in $\mathbf{c}$, the change is exactly $\delta\psi_2$ (shaded). Compact support confines it to its knot span; the RBF’s Gaussian tail still reaches the data interval (blue); a global polynomial changes the edge everywhere.

Table 8. Basis ablation (fixed Tsit5 solver, $G = 5$ grid points, 240 parameters, 10^4 epochs). κ_2 is the design-matrix condition number of Figure 12.

Basis function	Training MSE	Extrap. MSE	R^2	Rel. L_2	$\hat{L}$	κ_2	Wall-clock
Cubic B-spline	7.20×10^{-5}	8.38×10^{-5}	0.99997	0.32%	8.44	5.5	6680 s
Gaussian RBF	8.83×10^{-5}	8.92×10^{-5}	0.99997	0.33%	6.91	34.0	1712 s
Lagrange	1.07×10^{-4}	3.93×10^{-4}	0.99987	0.69%	6.70	3.7	6829 s
Chebyshev	9.80×10^{-5}	4.37×10^{-4}	0.99985	0.73%	6.48	2.5	2317 s
IQF	2.08×10^{-4}	9.61×10^{-4}	0.99967	1.08%	7.16	10.1	1824 s
RSWAF	1.00×10^{-4}	1.46×10^{-3}	0.99950	1.33%	8.72	7.9	1715 s
Newton	1.47×10^{-3}	3.65×10^{-2}	0.98746	6.66%	5.65	54.2	2187 s

and extrapolates $5 \times$ worse than RBF, the worst-conditioned of the working kernels. Locality matters more than conditioning once κ_2 is moderate. The two heavy-tailed kernels, IQF (algebraic $O(r^{-2})$ tails) and RSWAF, train as well as the polynomials but extrapolate worst among the working bases (11–16 $\times$ RBF). Newton is the outlier in both directions: it is the worst-conditioned basis and the only one that never reaches 10^{-3} training loss, and its lowest-in-table Lipschitz estimate (5.65) reflects underfitting, not desirable flatness. Figure 14 shows the same ranking in the phase plane.




Figure 14. Phase portraits on $[0, 14]$ for four of the seven bases (Tsit5). B-spline, RBF and Chebyshev all trace the true orbit closely at this scale; Chebyshev’s $5 \times$ larger extrapolation error is phase error rather than a visible departure. The Newton-basis model visibly misses the orbit near its upper and right-hand extremes.

3.4 Sensitivity to step size and observation noise

Two further one-factor sweeps vary the data rather than the model (Figure 15, Table 9). The **step-size sweep** changes the sampling interval Δt , and with it the number of training points and the solver step $h = \Delta t/2$. It is not monotone: $\Delta t = 0.1$ extrapolates best, $\Delta t = 0.05$ is $1.5 \times$ worse with twice the data and twice the cost, and $\Delta t = 0.2$ reaches the *lowest* training loss and the worst extrapolation ($5.1 \times$). The work-precision data exclude

truncation as the cause: Tsit5’s own error on the true field at these three steps is 2.6×10^{-18} , 4.7×10^{-14} and 2.2×10^{-10} in MSE, at least six orders of magnitude below any learned error. The coarse grid degrades extrapolation because 19 samples constrain the field less (it fits them more closely and generalises worse), and the fine grid because the same 10,000 epochs buy a similar training loss on a harder, longer unrolled graph. Step size is a data-density choice here, not a numerical-accuracy one.

Table 9. Step-size and observation-noise sweeps (Tsit5, RBF, 10^4 epochs). Noise is added only to the training window; all errors are measured against the clean truth.

Configuration	Training points	Training MSE	Extrap. MSE	Extrap. R^2	$\hat{L}$	Wall-clock
$\Delta t = 0.05$ ($h = 0.025$)	71	9.37×10^{-5}	1.34×10^{-4}	0.99995	7.46	3425 s
$\Delta t = 0.1$ (reference)	36	8.83×10^{-5}	8.92×10^{-5}	0.99997	6.91	1713 s
$\Delta t = 0.2$ ($h = 0.1$)	19	7.83×10^{-5}	4.51×10^{-4}	0.99985	6.64	882 s
$\sigma = 0.01$	36	1.23×10^{-4}	1.22×10^{-3}	0.99958	7.18	1686 s
$\sigma = 0.05$	36	7.41×10^{-4}	2.54×10^{-2}	0.99129	8.69	1672 s
$\sigma = 0.1$	36	2.28×10^{-3}	1.37×10^{-1}	0.95282	9.55	1534 s




Figure 15. Left: error against sampling interval, with $O(\Delta t)$ and $O(\Delta t^2)$ reference slopes; changing Δt also changes the number of training samples. Right: error against observation-noise level, with least-squares power-law fits; the dotted line is the noise variance σ^2 .

The **noise sweep** adds Gaussian noise of standard deviation σ to the 36 training observations. Across $\sigma \in \{0.01, 0.05, 0.1\}$ the extrapolation MSE follows a clean power law, $\text{MSE}_{\text{extrap}} \approx 12.9 \sigma^{2.02}$: extrapolation error is proportional to the noise variance, with a gain of about 13. The training-window error against the clean truth grows more slowly ($\propto \sigma^{1.24}$) and sits below σ^2 at every level ($0.23\sigma^2$ at $\sigma = 0.1$), so the ODE constraint does denoise the fit inside the window. What it cannot do is stop the residual field error from compounding into phase error beyond it, and the field roughens as it absorbs the noise ($\hat{L}$ rises from 6.91 to 9.55). Section 8 compares this sensitivity against sparse symbolic regression.

3.5 What the trained vector fields actually learned

Trajectory error measures one orbit. The saved checkpoints make it possible to interrogate the learned field $\mathbf{f}_\theta$ itself, as a dynamical system, with the invariants of Section 2 as ground truth.

Equilibria by Newton’s method. Solving $\mathbf{f}_\theta(\mathbf{u}^*) = \mathbf{0}$ by Newton-Raphson, $\mathbf{u} \leftarrow \mathbf{u} - J_\theta(\mathbf{u})^{-1}\mathbf{f}_\theta(\mathbf{u})$ with the Jacobian from automatic differentiation in float64, started at the true equilibrium $(3, 1.5)$, converges quadratically for every working model. For the Tsit5 model the residual norms are 2.1, 1.6×10^{-1} , 1.1×10^{-2} , 9.5×10^{-6} , 3.4×10^{-12} , then exactly zero: the residual is roughly squared at every step, the signature of quadratic convergence ($\|\mathbf{F}_{k+1}\|/\|\mathbf{F}_k\|^2$ stays between 0.04 and 0.43). The roots are far from the truth (Table 10): every RBF solver model places its equilibrium near $(1.9, 1.04)$, about 1.15 from $(3, 1.5)$, and it is an unstable

focus, with $\text{Re } \lambda \approx +1.2$ where the true system has $\text{Re } \lambda = 0$. The imaginary parts are close to the true 2.12. Chebyshev and Lagrange models have *stable* foci instead; Newton iteration on the Chebyshev model does not converge within 30 steps, and on the Newton-basis model it diverges.

Table 10. Numerical analysis of the trained fields. $\mathbf{u}^*$: learned equilibrium (Newton-Raphson from $(3, 1.5)$); λ : eigenvalues of $J_\theta(\mathbf{u}^*)$; period error from zero-crossings of the predicted prey signal over $[0, 28]$; H drift is $\max_t |H(\hat{\mathbf{u}}) - H_0|/H_0$ on $[0, 28]$. True values: $\mathbf{u}^* = (3, 1.5)$, $\lambda = \pm 2.121i$, $T = 3.3175$, drift 0.

Model	$\mathbf{u}^*$	$\ \mathbf{u}^* - \mathbf{u}_{\text{true}}^*\ $	$\text{Re } \lambda$	$ \text{Im } \lambda $	Period error	H drift
Euler	(1.80, 1.02)	1.30	+0.51	2.05	-1.7×10^{-3}	4.2%
Heun	(1.91, 1.03)	1.19	+1.11	2.03	-4.7×10^{-4}	3.1%
Midpoint	(1.93, 1.04)	1.16	+1.17	2.01	-3.5×10^{-6}	3.1%
RK4	(1.94, 1.05)	1.15	+1.21	1.99	$+7.8 \times 10^{-5}$	3.2%
DOPRI5	(1.94, 1.04)	1.16	+1.20	2.00	$+3.5 \times 10^{-5}$	3.2%
Tsit5 / RBF	(1.94, 1.04)	1.15	+1.20	2.01	$+1.7 \times 10^{-5}$	3.1%
B-spline	(1.66, 0.94)	1.46	+1.18	2.99	$+1.3 \times 10^{-4}$	1.3%
Lagrange	(2.67, 2.15)	0.73	-0.44	5.21	-9.9×10^{-4}	3.9%
Chebyshev	Newton iteration does not converge				$+6.0 \times 10^{-4}$	5.1%
IQF	(2.05, 1.19)	1.00	+0.26	1.70	-2.2×10^{-3}	76%
RSWAF	(2.15, 0.91)	1.03	+1.00	1.69	$+6.3 \times 10^{-4}$	5.4%
Newton	Newton iteration diverges				-6.4×10^{-3}	12%
MLP-ODE (SiLU), 10^4	(2.35, 1.12)	0.76	+0.30	2.75	$+8.0 \times 10^{-4}$	2.8%
MLP-ODE (tanh, paper), 10^4	(1.64, 0.79)	1.54	+0.79	2.43	$+3.1 \times 10^{-3}$	22.7%
KAN-ODE, 5×10^4	(1.51, 0.80)	1.65	+0.67	2.85	-1.6×10^{-4}	0.9%
MLP-ODE (SiLU), 5×10^4	(3.65, 1.30)	0.68	+0.95	1.83	$+1.0 \times 10^{-4}$	0.4%
MLP-ODE (tanh, paper), 5×10^4	(1.47, 0.72)	1.72	+1.09	3.16	$+6.8 \times 10^{-4}$	2.8%

The learned system is a limit-cycle oscillator. An unstable focus surrounded by an orbit that trajectories stay on is the signature of an attracting limit cycle, not of a conservative centre. Integrating each learned field from four initial conditions whose true orbits have $H \in \{0.85, 0.86, 2.0, 3.5\}$ confirms it (Figure 16): in the true system the four orbits are distinct nested curves, while under the KAN and SiLU fields – at 10^4 and at 5×10^4 epochs – all four trajectories converge onto the single training orbit, with late-time H between 2.000 and 2.008 for every starting point. The paper-settings tanh MLP behaves the same way at 5×10^4 epochs (late-time $H = 1.998$), but at 10^4 epochs it captures only three of the four starting points ($H = 2.032$); the trajectory from $(0.6, 0.6)$ diverges. Both architectures have replaced Lotka-Volterra’s one-parameter family of neutral orbits by one isolated, attracting cycle through the training data. This explains two things at once. The flat extrapolation error of the KAN-ODE (Table 15) is partly a property of the attractor: amplitude errors are damped back onto the cycle, so only phase error can accumulate, which is why the period error in Table 10 predicts the extrapolation ranking better than any field metric. And the models say nothing reliable about any other initial condition: they have learned an orbit, not the vector field.

Where the field is right. Figure 17 maps the field error $\|\mathbf{f}_\theta - \mathbf{f}\|$ over the phase plane. It is small only in a thin tube around the training orbit. On the orbit the error is 2.7–2.8% of the mean field magnitude for the 10k KAN and SiLU models and 0.6–1.1% at 50k, while the paper-settings tanh MLP is less accurate at 7.0% (10k) and 2.7% (50k); across the plotted box the median error is 20–34% for KAN and SiLU and 42–46% for the tanh MLP, and less than 3% of the box is below 1% error for any model. Longer training sharpens the tube without widening it: the 50k KAN’s median off-orbit error (34%) is higher than the 10k model’s (24%).

Invariant and error growth. Figure 18b tracks the first integral along each solver model’s own prediction. All of them hold H to within 3–4%, and the drift is oscillatory rather than secular, as it must be on an attracting cycle; the learned models’ invariant error is dominated by model error, orders of magnitude above the solver drift in panel (a) except for Euler. Figure 19 shows the pointwise error out to $t = 28$: the Midpoint and Tsit5




Figure 16. Trajectories of each learned field from four initial conditions (coloured dots) over $t \in [0, 40]$, against the four corresponding true orbits (grey). Red crosses mark the learned equilibria, black plus signs the true one. Columns: KAN-ODE, SiLU MLP-ODE, and the paper’s tanh MLP-ODE under the paper’s own learning rate and initialization; rows: 10^4 and 5×10^4 epochs. Every learned field attracts all four onto the training orbit except the 10^4 -epoch tanh MLP, under which the trajectory from $(0.6, 0.6)$ diverges.




Figure 17. Field error $\|\mathbf{f}_\theta(\mathbf{u}) - \mathbf{f}(\mathbf{u})\|$, normalised by the mean true field magnitude on the orbit (4.82), clipped to $[10^{-3}, 1]$. Thick black: training segment; blue: rest of the orbit; plus sign: true equilibrium. Columns and rows as in Figure 16.

models' error envelopes grow by $5.5\text{--}5.7 \times 10^{-4}$ per unit time, Heun's and Chebyshev's by $6\text{--}8 \times 10^{-3}$, and Euler's by 2.1×10^{-2} . Panel (b) previews Section 9: at 5×10^4 epochs the MLP error is lower than the KAN's throughout, while the 50k KAN has a *larger* far-horizon error than the 10k KAN (maximum 6.5×10^{-2} against 3.4×10^{-2} on $(14, 28]$), despite its better scored-window MSE.




Figure 18. Relative drift of the Lotka-Volterra first integral H (Equation 5). (a) The exact field integrated by each solver at $h = 0.05$: pure truncation drift (Euler's trajectory leaves the domain of H at $t = 19.4$). (b) Each trained model integrated by its training solver. Shaded: training window; dotted: end of the scored horizon.




Figure 19. Pointwise trajectory error on $[0, 28]$. (a) Solver and basis choice at 10^4 epochs (Midpoint and Tsit5 overlap). (b) KAN-ODE and MLP-ODE at 10^4 and 5×10^4 epochs.

What the edges look like. Figure 20 plots the twenty first-layer edge functions $\phi_{i,j}$ of the RBF and B-spline models over the range each input covers on the orbit. Over that range most edges are close to affine, with mild curvature concentrated at small populations where $\tanh$ is least saturated. None of them resembles a term of Equation 4; the bilinear products $u_1 u_2$ are assembled by composition through the second layer rather than carried by any single edge, consistent with the pruning result of Section 8.

3.6 Training dynamics across the sweeps

Figures 21 and 22 show the test MSE of every configuration during training, and Table 11 records when each one's training MSE crosses fixed thresholds. Five of the six solver curves lie on top of each other throughout; Euler separates from about epoch 2000 and stays above the others until the final approach, so solver choice changes where training ends rather than how it starts. From about epoch 6000 every solver shows periodic loss spikes of one to two orders of magnitude that recover within a few hundred epochs. The basis curves separate much earlier: B-spline is lowest almost throughout, Newton highest, and B-spline, Chebyshev and RSWAF reach 10^{-2} training MSE first (1900–2050 epochs, against RBF's 3112). Only B-spline converts that head start into better extrapolation.




Figure 20. First-layer edge functions (spline plus residual branch) of the trained RBF and B-spline models, each input varied over its range on the orbit with the other held at its mean. The edge with the largest range is highlighted.




Figure 21. Test MSE during training for the six solvers (RBF basis), logged every ten epochs; the dotted line is 10^{-4} .




Figure 22. Test MSE during training for the seven bases (Tsit5 solver), logged every ten epochs; the dotted line is 10^{-4} .

Table 11. First epoch at which the training MSE falls below each threshold (“–”: not reached within 10^4 epochs).

Solver	10^{-2}	10^{-3}	10^{-4}	Basis	10^{-2}	10^{-3}	10^{-4}
Euler	3533	6729	–	B-spline	1894	4826	8458
Heun	3144	5564	9512	Gaussian RBF	3112	5282	9617
Midpoint	3102	5285	9599	Chebyshev	1931	3892	9828
RK4	3153	5318	9798	Lagrange	2924	5407	–
DOPRI5	3132	5305	9796	IQF	2561	5726	–
Tsit5	3112	5282	9617	RSWAF	2047	4177	–
MLP-ODE (SiLU)	741	1297	9145	Newton	4577	–	–
MLP-ODE (tanh, paper)	3461	–	–				

What the ablation and the numerical analysis establish

- All six integrators satisfy their order conditions to round-off and show their design order; Tsit5’s leading error constant is $2.9\times$ smaller than DOPRI5’s.
- Above order two, solver choice does not change accuracy at $h = 0.05$; Midpoint is the cost optimum (Tsit5’s accuracy at 29% of the time).
- The network learns the field that its training solver needs, not the physical field: the Euler model is an accurate model of the inverse modified equation, and models do not transfer across solver families.
- Locality, not conditioning, separates the bases; on this system a third of the first layer’s grid is idle because states are positive.
- Every trained model, KAN or MLP, is an attracting limit-cycle oscillator around a spurious unstable equilibrium. It extrapolates the training orbit well and says nothing reliable about any other.

3.7 Diagnosing and fixing two cross-domain failures

Applying the Lotka-Volterra recipe unchanged to the damped pendulum and to SIR failed on both systems, each for its own reason. Both diagnoses follow the same discipline: measure a quantity about the model’s behavior *before* training rather than search the hyperparameter space after a failed run.

Pendulum: extending the training window. The pendulum recipe originally trained on $t \in [0, 3]$ and diverged in extrapolation ($R^2 = -1.318$). Extending the training window to $t_{\text{train}} = 5.0$ – a single flag – resolves it:

Question	Why does the Lotka-Volterra recipe, applied unchanged, diverge on the pendulum?
Expected (hypothesis)	SiLU’s residual branch cannot represent the sign-changing derivative $\dot{\theta} = \omega$ – an architectural limit.
Observed	Refuted. A longer training window fixes it under either activation; SiLU fits the sign change fine once the window is long enough.

Table 12. Pendulum fix: best configuration versus the original failure.

Configuration	Full-horizon MSE	Extrapolation R^2	Verdict
Original ($t_{\text{train}} = 3.0$)	8.99×10^{-1}	-1.318	diverges
Extended window ($t_{\text{train}} = 5.0$, SiLU)	4.48×10^{-2}	$+0.647$	adopted

An earlier hypothesis held that SiLU’s residual branch $W \cdot \text{silu}(x)$ cannot represent a sign-changing derivative such as $\dot{\theta} = \omega$. A factorial control comparing activation against training-window length refutes it directly: with the longer window, SiLU reaches an angular-velocity root-mean-square error of 0.0114 radians and reproduces the true turning point ($\theta_{\min} = -1.380$ against a true value of -1.385). The training window is what fixes the pendulum; the activation choice affects only how quickly the fit is reached, with the identity activation converging $\sim 1000\times$ faster at the short window without matching the longer window’s accuracy.




Figure 23. Fixed-pendulum trajectory reconstruction (left) and training/extrapolation loss curves (right), using the extended-window recipe.

SIR: a unit-scale mismatch. SIR diverged to an invalid (not-a-number) value at epoch 8686.

Question	What causes SIR’s divergence, and is it a scale problem, an optimizer defect, or a modeling gap?
Expected	Measuring the untrained network’s own field magnitude against the true dynamics, before any training, should reveal whether the two are mismatched in scale.
Observed	A $19\times$ speed mismatch at initialization. Rescaling time removes the divergence entirely; no optimizer or architecture change was needed.

Measuring the learned field magnitude *before any training step* explained the divergence directly. A Glorot-initialised KAN starts with $|f_{\theta}| \approx 0.203$ against SIR’s true field magnitude of 0.011 – a $19\times$ speed mismatch. To see why this is catastrophic rather than merely inaccurate, consider that backpropagation through an explicit solver differentiates roughly $T/\Delta t$ nested nonlinear steps, and the resulting gradient sensitivity grows exponentially as e^{LT} , where L is the Lipschitz constant of $\mathbf{f}_{\theta}$ and T is the integration horizon. Lotka-Volterra trains at

$T = 3.5$ with $|f| \sim 10^0$ and converges; the pendulum trains at $T = 5.0$ with $|f| \sim 10^0$ and converges after the window fix. SIR sits at $T = 50$ with $|f| \sim 10^{-2}$ – the product LT is large enough that the gradient overflow recorded at epoch 11, not the later epoch 8686, is where training actually fails.

The optimizer was behaving rationally: with an epoch-0 loss of 3.16×10^2 , collapsing the field to $\mathbf{f}_\theta \rightarrow \mathbf{0}$ reduces the loss to $\approx 10^{-1}$ – a $3,000\times$ improvement reachable immediately. SIR’s natural timescale is the recovery period $1/\gamma = 10$ days; substituting $\tau = t/10$ gives $\mathbf{dy}/d\tau = 10\mathbf{f}(\mathbf{y})$ over $\tau \in [0, 5]$, landing SIR in the same regime the other two systems already train in. Projecting the learned field onto the zero-sum subspace $f_{\text{proj}}(\mathbf{y}) = f(\mathbf{y}) - \frac{1}{n} \sum_i f_i(\mathbf{y})$ fixes a separate mass-drift symptom by enforcing the SIR invariant (Equation 6) as a property of the flow rather than a penalty on the fit. A third fix, a *vanishing-field gate*, multiplies the learned field by the infected fraction I , so that $I = 0$ is a manifold of equilibria exactly as in Equation 6, where every term carries a factor of I . Unlike the first two, this fix assumes something about the physics, so Table 13 reports the first two on their own and then with the gate added. Rescaling and projection alone remove the divergence and cut extrapolation RMSE by $4.4\times$; the gate gives a further $25\times$, and only the full recipe reaches a high extrapolation R^2 . Over this window the true signal varies very little (standard deviation 0.002–0.017), so R^2 is harsh and RMSE is the more informative number.

Table 13. SIR before and after the fixes, at the full 10,000-epoch budget. “Rescale + project” applies time rescaling and the conservation projection, which assume nothing beyond Equation 6; “+ gate” adds the vanishing-field gate. Mass error is $\max_t |S + I + R - 1|$.

Metric	Before fix	Rescale + project	+ gate
Training MSE	4.39×10^{-4}	1.08×10^{-5}	2.28×10^{-8}
Extrapolation RMSE	0.0614	0.0140	5.7×10^{-4}
Mass error	0.0045	0.0000	0.0000
Extrapolation R^2	−20.97	−0.145	+0.9981
Maximum gradient norm	5.82×10^{18}	92.3	1.19
Final model	invalid (NaN at epoch 8686)	finite	finite




Figure 24. SIR with all three fixes: reconstruction of S and I over the full horizon (left; the plot’s generic legend labels them prey and predator) tracks the true dynamics closely; training and extrapolation loss curves (right) decrease together without the gradient overflow that terminated the original run.

Other fixes considered. Several additional interventions were tested and are recorded for completeness rather than adopted:

Methodological lesson carried into every later study

Both diagnoses came from measuring the learned field’s own magnitude at initialization, before training ever ran. This init-time-measurement discipline reappears verbatim in the epidemic-curve test of Section 8.

Table 14. Every stability intervention tested, including the rejected ones.

Intervention	Verdict	Evidence
Longer training window (pendulum)	Adopted	R^2 improves from -1.318 to $+0.647$; works for both activations tested
Non-finite gradient guard with automatic abort	Adopted	Caught the divergence of an intermediate SIR recipe (conservation penalty plus loss weighting) at epoch 3055 and preserved a valid checkpoint instead of overwriting it with an invalid one
Time rescaling (SIR)	Adopted	Removes both the divergence and the frozen-field behavior
Conservation via subspace projection (SIR)	Adopted	Mass error reduced from 4.5×10^{-3} to 0.0000 (exact by construction)
Vanishing-field gate (SIR)	Adopted	Extrapolation RMSE reduced a further $25\times$ ($0.0140 \rightarrow 5.7 \times 10^{-4}$); the only fix that assumes a physical property
Identity activation (pendulum)	Speed only	Roughly $1000\times$ faster to fit at the short window, but worse extrapolation and stability at the long window
Standard-deviation loss weighting	System-dependent	Worst-performing pendulum option; superseded for SIR by time rescaling
Conservation via penalty term	Rejected	Wrong mechanism – the training window’s mass was already close to conserved, so the penalty had nothing to correct (mass error still 0.0086)
Widened basis-function grid limits	Rejected	Performed worse than the untouched baseline configuration

3.8 Extended extrapolation: the KAN advantage widens with horizon

With the fixes in place, the Lotka-Volterra extrapolation horizon was extended from four periods ($t \in (3.5, 14]$) to eight periods ($t \in (3.5, 28]$) using the same trained model, with no retraining.

Table 15. Extrapolation MSE as the horizon doubles, same trained models.

Model	Training $[0, 3.5]$	Near extrapolation $(3.5, 14]$	Far extrapolation $(14, 28]$	Degradation
KAN-ODE (RBF, 240 parameters)	8.83×10^{-5}	8.92×10^{-5}	9.72×10^{-5}	$1.1\times$
MLP-ODE (SiLU, 252 parameters)	9.50×10^{-5}	4.78×10^{-4}	3.95×10^{-3}	$8.3\times$
MLP-ODE (tanh, paper lr/init, 252 parameters)	1.11×10^{-3}	9.65×10^{-3}	5.95×10^{-2}	$6.2\times$

The KAN-ODE’s error is essentially flat across eight periods ($8.83 \rightarrow 8.92 \rightarrow 9.72 \times 10^{-5}$, a 10% rise from the training window to double the original horizon), while the parameter-matched MLP degrades $8.3\times$ over the same extension. The KAN-ODE advantage therefore widens with horizon: from $5.4\times$ at four periods to **$40.6\times$** at eight. This is a materially stronger claim than the near-horizon comparison already presented: KAN-ODE error does not accumulate, so the gap grows with every additional period.

3.9 Pendulum energy dissipation: an independent physical check

The fixed pendulum recipe captures 94.2% of the true energy dissipation and tracks $E(t)$ to 0.6% inside the training window, but the strict monotonic-decay invariant ($dE/dt \leq 0$, Equation 7) is violated on 40 of 200 predicted steps. The model is under-damped: the true angular displacement decays to 0.004 by $t = 9.1$ while the predicted trajectory continues oscillating at 0.237. This energy diagnostic is an independent confirmation of the under-damping, expressed as a violated conservation law rather than a trajectory RMSE (Figure 25).




Figure 25. Fixed pendulum ($\mu = 0.5$, $t_{\text{train}} = 5$): phase portrait (left) and total mechanical energy $E(\theta, \omega)$ (right). Inside the training window the predicted energy follows the true curve closely; past the split at $t = 5$ it stays above the truth and rises on 40 of 200 predicted steps, the under-damping described in the text.

4 Adjoint Sensitivity vs. Direct Autograd: Memory and Wall-Clock Cost

Two methods exist for computing gradients through a neural ODE. **Direct autograd** treats the numerical integrator as an ordinary computational graph: each of the solver’s N_t forward steps is recorded, and the reverse pass walks back through every recorded step in turn, so peak memory grows linearly with the number of steps taken. **Adjoint sensitivity** [3] instead treats the forward solve as a black box and derives a second, augmented ODE that produces the same gradient without storing every intermediate state. Writing $\mathbf{a}(t) = \partial L / \partial \mathbf{u}(t)$ for the adjoint state, it satisfies its own backward-time differential equation,

$$\frac{d\mathbf{a}(t)}{dt} = -\mathbf{a}(t)^\top \frac{\partial \mathbf{f}_\theta}{\partial \mathbf{u}}(\mathbf{u}(t)), \quad \frac{dL}{d\theta} = \int_{t_1}^{t_0} \mathbf{a}(t)^\top \frac{\partial \mathbf{f}_\theta}{\partial \theta}(\mathbf{u}(t)) dt, \quad (18)$$

solved backward from t_1 to t_0 alongside the original state $\mathbf{u}(t)$ (recomputed rather than stored), so memory depends only on the size of the state and parameter vectors, not on how many steps the forward solve took – an $O(1)$ memory cost in trajectory length, against direct autograd’s $O(N_t)$. This project’s own implementation documentation adopts direct autograd on the stated grounds that this $O(1)$ advantage does not matter for this project’s short trajectories ($N_t \in [36, 101]$). This section measures that claim directly on an RTX 3060, holding the integrator fixed (classical RK4 at matched step size on both paths; a forward-trajectory check confirms the two implementations agree to $\sim 10^{-7}$ relative error) so that only the differentiation method varies.





Figure 26. Direct autograd (top) keeps every intermediate state u_i in memory so the backward pass can reuse it, giving $O(N_t)$ memory. Adjoint sensitivity (bottom) instead solves the backward-time equation of Equation 18, recomputing $u(t)$ as needed rather than storing it, giving $O(1)$ memory regardless of N_t .

Question	Does adjoint sensitivity’s constant memory cost make it the better choice at this project’s trajectory lengths, as the project’s own documentation assumes?
Expected	Direct autograd should be cheaper in memory at these short lengths, with adjoint’s advantage only appearing once trajectories grow much longer.
Observed	Adjoint already uses less memory at every length tested. The deciding cost is wall-clock, not memory: adjoint is $1.9\text{--}2.1\times$ slower.

Table 16. Peak VRAM and wall-clock, direct autograd vs. adjoint sensitivity.

System	Trajectory length N_t	Method	Peak VRAM (MB)	Seconds per 1000 epochs
Lotka-Volterra	36	Direct autograd	18.61	317.5
Lotka-Volterra	36	Adjoint	17.08	672.9
SIR	101	Direct autograd	21.52	984.8
SIR	101	Adjoint	17.12	1871.6

**Figure 27.** Peak memory vs. trajectory length $N_t \in \{36, \dots, 3201\}$: direct autograd grows linearly (≈ 0.045 MB per output point); adjoint stays flat (0.7% increase over an $89\times$ length increase).

Finding. The measured memory data contradicts the stated premise behind the original choice, while still supporting its conclusion. The $O(N_t)$ vs. $O(1)$ scaling in Equation 18 predicts a crossover length below which direct autograd’s smaller constant factor should still win on raw memory, since it avoids the adjoint method’s own bookkeeping overhead; the data shows no such crossover inside the tested range. Adjoint uses *less* memory than direct autograd at every tested length, including the shortest ($N_t = 36$: 17.08 vs. 18.61 MB) – Figure 27 shows direct autograd’s per-step storage cost is already large enough at this project’s scale to outweigh whatever constant overhead the adjoint method’s own bookkeeping carries. Extrapolating the measured linear slope of direct autograd’s memory growth, N_t would need to reach $\approx 225,000$ before it threatens a 12 GB card. The deciding factor for this project is wall-clock: solving the continuous adjoint ODE backward in time costs a consistent $1.9\text{--}2.1\times$ overhead at both tested lengths, turning a 53-minute Lotka-Volterra run into ~ 112 minutes for a memory saving that never becomes load-bearing (peak usage is 0.17% of card capacity). Gradient relative error is 2.85×10^{-7} and 4.88×10^{-7} respectively, well inside a 5×10^{-3} tolerance – both methods agree numerically, so the choice reduces to a resource trade-off. Direct autograd is the faster option for every dataset in this project’s scope ($N_t \leq 201$).

5 Gradient-Norm Dynamics and Solver Order

The benchmark ablation found Forward Euler uniquely unable to reach 10^{-4} training loss within budget, with $28\times$ worse extrapolation MSE than Tsit5 (Table 7). The candidate mechanism tested here: a first-order solver commits $O(\Delta t)$ local error every step, which could inject measurably noisier gradients and destabilize Adam’s moment estimates.

Question	Does a lower-order solver’s larger truncation error show up as measurably noisier gradients during training?
Expected	Gradient-norm noise should fall as solver order rises, correlating with the accuracy ranking already seen in the benchmark ablation.
Observed	No reliable correlation ($\rho = -0.088$, $p = 0.87$); the ranking is non-monotone and its sign flips depending on an arbitrary warm-up cut. Refuted.

The noise measure is the standard deviation of epoch-to-epoch *relative* gradient change, scale-invariant by construction,

$$\eta = \text{std}(\Delta \log_{10} g), \quad \Delta \log_{10} g_i = \log_{10} g_{i+1} - \log_{10} g_i, \quad (19)$$

compared against solver order p by Spearman rank correlation rather than the more familiar Pearson correlation. Pearson correlation measures how linearly two quantities move together, which presumes solver order is a quantity whose actual numeric spacing is meaningful – that the jump from order 1 to order 2 is comparable in size to the jump from order 4 to order 5. There is no reason to expect η to depend on solver order that linearly; the hypothesis under test is only that noise falls *monotonically* as order rises, which is exactly what Spearman correlation measures, by comparing ranks rather than raw values. It also handles the deliberate tie at $p = 2$ (Heun and Midpoint are both second-order methods) natively, where Pearson would need an arbitrary tie-breaking rule.

Table 17. Gradient-noise measure η vs. solver order, six solvers, $n = 6$.

Solver	Order p	Noise η	Best training MSE	Noise rank
Forward Euler	1	0.2805	1.377×10^{-4}	1 (noisiest)
Classical RK4	4	0.2726	9.413×10^{-5}	2
Tsit5	5	0.2683	8.833×10^{-5}	3
DOPRI5	5	0.2660	9.283×10^{-5}	4
Midpoint RK2	2	0.2634	8.857×10^{-5}	5
Heun RK2	2	0.2621	8.315×10^{-5}	6

Spearman $\rho = -0.088$, $p = 0.87$ ($n = 6$). Three findings jointly refute the hypothesis: the effect size is a 7% spread across all six solvers; the ranking is non-monotone (RK4, order 4, is the second-noisiest solver); and the sign of ρ flips across warm-up cuts ($-0.088 \rightarrow -0.441 \rightarrow +0.706$ as the first 0, 2000, 5000 epochs are dropped), so no cut reaches significance and the direction of any relationship depends on an arbitrary analyst choice.

Finding. Gradient-norm roughness does not explain Euler’s convergence floor. All six traces are instead dominated by shared, order-independent spike bursts (1176 spike-epochs for Euler, 866 for Tsit5 in the last 4000 epochs, near-identical burst structure). Four checks rule out a numerical defect: zero non-finite gradients or losses across 60,000 logged epochs; loss decreases monotonically in its running minimum for every solver; loss is lower five epochs after a spike in 64.5% of cases; and median loss is elevated on spike epochs (3.90×10^{-4} vs. 1.80×10^{-4} on calm epochs), ruling out a spurious-gradient/still-loss signature. The most plausible reading is edge-of-stability optimizer behaviour: gradient descent with a fixed step size and no learning-rate schedule is known to settle into a regime where it repeatedly overshoots the walls of a narrowing loss valley as training approaches a sharp minimum, producing exactly this pattern of a large transient spike followed by fast recovery, rather than either smooth convergence or a genuine divergence. Because this behavior appears with near-identical structure in Euler’s and Tsit5’s traces alike, it is a property of the optimizer and the loss surface’s local geometry, not of the integrator’s truncation error. Euler’s convergence floor is better explained by the earlier




Figure 28. Gradient-norm traces, all six solvers: large sustained spike bursts (one to two orders of magnitude above the local median) dominate every trace, including Tsit5, independent of solver order.

systematic-bias hypothesis (Section 3): a first-order integrator's $O(\Delta t)$ local error biases *where* the optimizer settles, without making the path it takes to get there measurably rougher.

6 Learnable Basis-Function Blending

The benchmark ablation found B-spline (compact local support, $3.9\times$ RBF’s cost) and RBF (smooth global support, cheap) statistically tied on accuracy. Rather than choosing one basis function $B_i(x)$ in Equation 2 in advance, this study makes the choice itself a trained quantity: a learnable convex blend of the two candidate bases inside every edge,

$$\phi(x) = \alpha \cdot \text{Spline}(x) + \beta \cdot \text{RBF}(x), \quad (\alpha, \beta) = \text{softmax}(\text{logits}), \quad (20)$$

trained end-to-end from a neutral start ($\alpha = \beta = 0.5$) with no hand-set schedule, on Lotka-Volterra and the stability-fixed pendulum (Section 3). The blend weights are produced by a softmax over two learnable logits rather than, say, a single sigmoid-gated weight, because a softmax guarantees $\alpha + \beta = 1$ by construction – the two bases always form a genuine convex combination, so the gate can only trade one basis off against the other rather than scaling the overall edge output up or down as a side effect. A single pair of logits is shared across every edge in the network, so the whole model learns one global blend rather than letting each edge pick its own; this was a deliberate simplification to keep the number of new parameters at 2, not a claim that per-edge blending would not also work. If a single global blend systematically beats both pure bases, it would suggest the basis-function choice in Equation 2 is itself worth learning rather than fixing by hand, for every future KAN-ODE application.




Figure 29. The hybrid-basis edge of Equation 20: every edge evaluates both bases in parallel, and a shared softmax gate produces the scalar weights α, β that multiply each basis’s output before they are summed into $\phi(x)$ – the choice of basis is a trained quantity, not a decision fixed at construction time.

Question	Can a network do better by learning its own blend of B-spline and RBF, rather than committing to one basis in advance?
Expected	A learned blend should match or beat the better pure basis, by adaptively favoring local support where useful and global support elsewhere.
Observed	Hybrid never beats the better pure basis on either system. On the pendulum, it actively overfits past 3000 epochs in a way pure RBF does not.

Table 18. Hybrid basis vs. pure bases, full 10,000-epoch budget.

System / basis	Training MSE	Extrapolation MSE	R^2	Final gate
Lotka-Volterra – RBF (pure)	8.83×10^{-5}	8.92×10^{-5}	1.0000	–
Lotka-Volterra – B-spline (pure)	7.20×10^{-5}	8.38×10^{-5}	1.0000	–
Lotka-Volterra – Hybrid	1.47×10^{-4}	1.78×10^{-4}	0.9999	$\alpha=0.12, \beta=0.88$
Pendulum – RBF (pure)	9.28×10^{-5}	0.090	+0.647	–
Pendulum – Hybrid (10,000 epochs)	9.44×10^{-5}	0.809	–2.17	$\alpha=0.02, \beta=0.98$

Lotka-Volterra: no benefit, real cost. The gate settles into a genuine partial blend, retaining $\approx 12\%$ B-spline, which confirms the mechanism works as designed. Hybrid ties both pure bases on final accuracy while paying roughly the *sum* of their per-epoch compute.




Figure 30. Learned α/β gate evolution: Lotka-Volterra settles into a genuine partial blend (left); the pendulum gate reaches $\approx 99.5\%$ RBF by epoch 250 and ends at 98% (right); it settles well before the overfitting onset discussed below.

Pendulum: overfitting emerges only past a specific budget. A short probe run at 2000 epochs first surfaced a practical wrinkle: the gate moved in the correct direction (toward RBF) but too slowly, so the network still carried meaningful B-spline contamination at an epoch count where the pure-RBF pendulum recipe was already nearly converged. Giving the gate its own, faster learning rate ($15\times$ the base rate, direction untouched) closed nearly all of that gap at the same probe budget (Figure 31, left) and was carried into the full run.




Figure 31. Left: the 2000-epoch probe run’s gate trace, showing the slow drift toward RBF that motivated the gate-learning-rate fix. Right: the gate trace for the 3000-epoch pre-overfitting-optimum run discussed below, settling to $\beta \approx 0.996$.

At the full 10,000-epoch budget, hybrid matches pure RBF’s training fit almost exactly, but its extrapolation collapses ($R^2 = -2.17$): the monitor loss bottoms out near epoch 2950 (≈ 0.054) then rises to a ≈ 0.4 plateau for the remaining 7000 epochs while training loss keeps improving – textbook overfitting, visible directly in the saved training history. The gate settles well before this onset: β reaches ≈ 0.995 by epoch 250, some 2750 epochs before overfitting begins, and pure RBF trained on the identical recipe shows no such pattern at all. Capping training at the pre-overfit optimum (epoch 3000, chosen from the 10,000-epoch run’s own point of minimum monitor loss) shows hybrid *ahead* of pure RBF at the same epoch count (0.055 vs. 0.081 full-horizon MSE) with extrapolation $R^2 = +0.580$, approaching pure RBF’s fully-converged number ($+0.647$). Figure 32 shows both runs side by side: the loss curves of the full run make the overfitting onset visible, and its phase portrait shows the extrapolation collapsing into a small wrong loop, while the 3000-epoch run still follows the decaying spiral.

Verdict

The learnable hybrid basis never outperforms the better pure basis on either system at any budget tested. Its gate mechanism works exactly as designed – gradient-driven, converging differently and sensibly per system – but that alone does not buy an accuracy edge, and on the pendulum specifically a network that *starts* as a blend generalizes measurably worse over long training than one that starts pure RBF, despite an almost-identical final gate and matching training accuracy.




Figure 32. Hybrid basis on the pendulum. Top: the full 10,000-epoch run; training loss keeps falling while the extrapolation (monitor) loss bottoms out near epoch 2950 and climbs back to ≈ 0.4 (left), and the extrapolated orbit collapses into a small wrong loop (right). Bottom: the same recipe stopped at 3000 epochs; train and monitor loss fall together (left) and the extrapolation follows the decaying spiral (right).

7 Stiffness–Solver Stability and Loss-Landscape Analysis

Numerical stiffness – a system whose fastest and slowest timescales differ by orders of magnitude – is where explicit integrators are classically expected to struggle. This study sweeps damping ratio $\mu \in \{0.1, 0.5, 1.0, 2.0, 5.0, 8.0\}$ against four explicit solvers {Euler, Midpoint, RK4, Tsit5} [9] on the damped pendulum, at a fixed step $\Delta t = 0.05$, to map which (solver, μ) combinations converge and which fail, and to test whether any failure is actually a stiffness effect.

Question	As damping ratio increases the pendulum’s dynamics stiffen – do lower-order explicit solvers genuinely fail as stiffness grows?
Expected	Higher damping ratios should push some (solver, μ) combinations outside that solver’s classical stability region, producing divergence.
Observed	Linear stability theory predicts zero divergences at every tested μ , and the sweep confirms it. Every observed failure traces to random seed, not stiffness.

7.1 The stability question resolves quickly

Linearizing at equilibrium gives eigenvalues $\lambda = (-\mu \pm \sqrt{\mu^2 - 4g/L})/2$. For every underdamped $\mu < 2\sqrt{g/L} \approx 6.26$, $|\lambda| = \sqrt{g/L} \approx 3.13$ is constant, so $h|\lambda| \approx 0.157$ at $\Delta t = 0.05$ – far inside even Euler’s stability region $h|\lambda| \leq 2$ [4]. Even the overdamped case ($\mu = 8.0$, fast eigenvalue $\lambda_{\text{fast}} \approx -7.7$) gives $h|\lambda_{\text{fast}}| \approx 0.385$, still comfortably stable. This predicts, and the sweep confirms, zero divergences across all 24 cells: not one run produced an invalid gradient at any point in training. All four solvers tested are explicit Runge-Kutta methods with a bounded stability region; no implicit solver is implemented in this project, so this analysis cannot speak to whether an implicit method would behave differently at a genuinely stiffer configuration.

7.2 What actually drives the failures: initialization

Table 19. Original 24-cell sweep (single seed 42), best training MSE and verdict.

Damping ratio μ	Euler	Midpoint	RK4	Tsit5
0.1	unstable, 1.7×10^{-2}	unstable, 3.0×10^{-2}	unstable, 2.2×10^{-2}	unstable, 2.4×10^{-2}
0.5	converged , 4.6×10^{-3}	converged , 7.9×10^{-3}	converged , 4.7×10^{-3}	converged , 9.2×10^{-3}
1.0	unstable, 1.0×10^{-1}	converged, 5.4×10^{-4}	converged, 5.3×10^{-4}	converged, 4.0×10^{-4}
2.0	unstable, 1.8×10^{-2}	unstable, 1.8×10^{-2}	unstable, 1.8×10^{-2}	unstable, 1.8×10^{-2}
5.0	converged, 6.1×10^{-5}	converged, 7.0×10^{-5}	converged, 6.9×10^{-5}	converged, 6.8×10^{-5}
8.0	converged, 2.1×10^{-5}	converged, 2.3×10^{-5}	converged, 2.6×10^{-5}	converged, 2.6×10^{-5}

Two anomalies stand out: $\mu = 2.0$ fails identically across all four solvers to two significant figures, and $\mu = 1.0$ fails only for Euler. Retraining Euler at $\mu = 2.0$ with two more seeds isolates the cause (seed 42: unstable, 1.82×10^{-2} ; seed 1: unstable, 1.68×10^{-2} ; seed 7: **converged**, 7.09×10^{-5}) – the trap depends on the random initialization. Extending this to all four solvers $\times$ six $\mu \times$ three seeds (72 cells, checkpoints saved throughout) confirms the pattern generally: only the two extremes ($\mu = 0.1$, uniformly hard; $\mu \in \{5.0, 8.0\}$, uniformly easy) hold their verdict across every seed. Every transitional $\mu \in \{0.5, 1.0, 2.0\}$ has at least one of three seeds land in a bad local minimum, and at $\mu = 1.0$ two *different* seeds fail for two different reasons: seed 1 produces a universally bad initialization shared by all four solvers (MSE $\approx 9.7\text{--}9.8 \times 10^{-2}$, the same signature as the $\mu = 2.0$ shared trap), while seed 42 carries a second, Euler-specific failure on top of it. Euler’s apparent $\mu = 1.0$ weakness in Table 19 is therefore the sum of two separate effects, one shared and one solver-specific, compressed into a single misleading number.

Consequence for the original heatmap

The 24-cell sweep used one shared seed for all four solvers. All four solvers failing identically at $\mu = 2.0$ reflects four solvers inheriting the same coin flip from a shared initialization, at this step size the seed $\times$ μ interaction determines outcome far more than solver choice does.




Figure 33. Stability verdict heatmap, 24-cell single-seed probe sweep, $\Delta t = 0.05$.

7.3 The mechanism: a genuine spurious equilibrium

A checkpoint saved for Euler’s trapped $\mu = 2.0$ run, evaluated at the true equilibrium and at a nearby grid-search minimum, shows the learned field f_θ has a large nonzero derivative at $(0, 0)$ ($|f_\theta| = 2.74$) and a genuine, stable, spurious equilibrium 0.05–0.07 units away ($|f_\theta| = 0.0086$). This holds independently for all four solvers’ own trained models ($|f_\theta(0, 0)| \in [2.65, 2.88]$ across euler, midpoint, rk4, tsit5): all four independently-trained networks converge to essentially the same wrong fixed point.

The full predicted trajectories tell the same story over time rather than in a single snapshot. Figure 35 shows every solver’s predicted angle, angular velocity, and error at $\mu = 2.0$ settling onto the same wrong, generic trajectory; Figure 36 shows the $\mu = 1.0$ case, where only Euler’s predicted angle visibly departs from the true one.

7.4 Weight-space loss landscape: an independent measurement

The phase-space plots above describe the learned dynamics; a separate question is how forgiving the *training loss* is around each solution – whether a small nudge to the weights leaves the fit essentially unchanged or makes it collapse. This was measured by stepping each trained model’s weights along two random directions and re-evaluating training error on a grid, following the visualization technique of Li et al. [7] adapted to KAN-ODE’s 360-parameter weight space. Every cell traces the same qualitative shape – a dip at the trained point that rises moving outward – and is summarized by its “rise”: the number of orders of magnitude the log-loss climbs by the edge of the grid.

The converged models (rise ≈ 9) sit at the bottom of a narrow, steep basin: stepping even a small distance from the trained weights in either random direction drives the log-loss up by nine orders of magnitude, so the surrounding terrain looks like a sharp valley with the solution pinned at its floor. The trapped cells (bold, rise ≈ 4 –5.5) sit in a basin roughly half as steep – the terrain around them is closer to a broad, gently sloped plain than a valley, and the same size step barely raises the loss by comparison. This is the opposite of the usual intuition that a flat minimum generalizes better: here, the flat region is where training settled for a generic,




Figure 34. Learned vector-field landscape $|f_\theta(\theta, \omega)|$, per solver's own trained model. Left ($\mu = 2.0$): every solver's true equilibrium (star) sits partway up a slope, while a separate, lower point nearby is where each trajectory actually settles. Right ($\mu = 1.0$): Euler's star sits high on a ridge, well apart from where its trajectory spirals down and rests; for the other three solvers, star and resting point coincide at the valley floor.




Figure 35. $\mu = 2.0$, all four solvers: predicted angle, angular velocity, and error over time, alongside each solver's phase portrait.




Figure 36. $\mu = 1.0$, all four solvers: the same panel layout, showing Euler's trajectory diverging from the true angle while the other three track it closely.

Table 20. Loss-landscape rise, single direction-pair (seed 0).

μ	euler	midpoint	rk4	tsit5
1.0	5.53	9.03	8.99	9.12
2.0	4.23	4.24	4.22	4.20
5.0	8.75	8.75	8.76	8.76
8.0	9.14	9.15	9.12	9.12

low-effort answer, and the steep, narrow basin is where it actually learned the pendulum’s dynamics.

Figure 37 renders this directly as a surface: the terrain around Euler’s trapped $\mu = 2.0$ solution is a broad, gently rolling plain, with no sharply defined bottom, in clear visual contrast to Figure 39’s converged $\mu = 1.0$ solution, which sits at the bottom of a narrow, sharply walled valley – the same contrast the rise numbers above report as a factor of roughly two, made visible. Figure 38 shows the trapped $\mu = 1.0$ case under the same solver whose $\mu = 2.0$ terrain is shown alongside it, confirming the flat terrain is a property of the trapped solution rather than of any one damping value.




Figure 37. Euler, $\mu = 2.0$ (trapped): the training-loss surface around the trained weights, viewed under three independent random direction-pairs. The basin near the trained point (the marked star) is broad and shallow in all three views; the third panel’s steep spike on the far side is the outlier direction discussed below, not the basin itself.




Figure 38. Euler, $\mu = 1.0$ (also trapped): the same broad, shallow basin shape recurs at a different damping value under the same solver.

Robustness check across three direction-pairs. Repeating the measurement with two more independent random direction pairs per cell (72 measurements total) found one direction-pair acting as a shared outlier across nearly every cell. Because every model shares the same architecture, one particular raw random draw runs toward a steeper wall in all 24 cells at once, inflating that direction’s rise by roughly $2\text{--}3\times$ everywhere. The size of the measured effect depends on how many direction choices are averaged: the single-seed $\approx 2\times$ gap between trapped and converged cells narrows to roughly 20–30% once all three direction-pairs are averaged, since the outlier direction inflates the trapped cells’ mean more than the converged cells’ mean. The ranking




Figure 39. Tsit5, $\mu = 1.0$ (converged): the basin around the trained point is visibly narrower and drops noticeably lower, in every one of the three direction-pairs, than either trapped case above.

between cells holds regardless; only the apparent size of the gap changes. The multi-seed average is the figure to trust; the single-seed table above illustrates the shape of the effect at a scale that is easier to see.




Figure 40. Mean rise per (solver, μ), 3-seed average with seed-to-seed range as error bars, colored by verdict.

The cells trapped by the spurious-equilibrium mechanism – $\mu = 2.0$ across all solvers, and Euler’s own $\mu = 1.0$ – sit around 20–30% lower on this chart than converged cells, a modest but measured gap rather than an order-of-magnitude separation. $\mu = 0.1$ carries the “unstable” label for an unrelated reason: its trajectory never settles within the training window at all, and its bars sit in the same range as the converged $\mu \in \{5.0, 8.0\}$ cells, occasionally above them. The flatness effect belongs specifically to the spurious-equilibrium trap cases, and does not generalize to every cell the pipeline labels “unstable.”

Verdict

Every divergence-free prediction from linear stability theory held exactly, so nothing in this sweep is a classical numerical-stiffness failure. The mechanism at work is a seed-dependent local minimum in a 360-parameter optimization landscape, producing a genuine, solver-independent spurious equilibrium in the learned dynamics. Three independent lines of evidence converge on this: per-solver trained-weight evaluation, phase-space vector-field visualization, and an independently-measured, direction-robust weight-space loss landscape.




Figure 41. Converged fraction (left) and median training MSE (right) over 3 seeds, all 4 solvers $\times$ 6 μ (72 cells). The two agree everywhere: low converged-fraction cells are exactly the cells with elevated median error.

8 Sparse Regression Comparison and Outbreak-Curve Validation

This section compares two fundamentally different ways of discovering governing equations from data – sparse symbolic regression and gradient-based neural fitting – and separately tests whether the cross-domain stability fixes of Section 3 generalize beyond the mechanistic system that produced them.

8.1 Noise robustness: KAN-ODE vs. sparse symbolic regression

Sparse Identification of Nonlinear Dynamics (SINDy) [2] fits a governing equation as a sparse linear combination of candidate terms drawn from a fixed library – here, every polynomial in x and y up to degree 2, i.e. $\{1, x, y, x^2, xy, y^2\}$. Coefficients are estimated by sequentially thresholded least squares: an ordinary least-squares fit is performed against numerically estimated derivatives, every coefficient below a fixed threshold is zeroed out, and the fit is repeated on the surviving terms until the active set stops changing. This is what makes the recovered model sparse and, in principle, human-readable – unlike the KAN’s 240 free parameters, SINDy returns a short list of surviving terms with numeric coefficients attached, directly comparable to the true equation’s own coefficients. The coefficient threshold is held fixed across the entire noise sweep so that the comparison measures noise sensitivity rather than per-level tuning. It is fit against the same noise-corrupted Lotka-Volterra observations the earlier KAN checkpoints (Section 3) were trained on, scored against clean ground truth.

Question	How does a sparse, symbolic method compare to a dense, 240-parameter neural method as observation noise increases?
Expected	Sparse regression’s tiny library should make it more robust to noise; KAN-ODE’s larger parameter count should make it more accurate.
Observed	Both expectations hold, but the KAN-ODE advantage collapses 170 $\times$ across the sweep, from 1110 $\times$ at zero noise to 6.5 $\times$ at $\sigma = 0.1$.

Table 21. Extrapolation MSE vs. observation noise σ .

Noise level σ	Sparse regression	KAN-ODE	KAN-ODE advantage
0.00	9.90×10^{-2}	8.92×10^{-5}	1110 $\times$
0.01	1.58×10^{-1}	1.22×10^{-3}	130 $\times$
0.05	1.48×10^{-1}	2.54×10^{-2}	5.8 $\times$
0.10	8.89×10^{-1}	1.37×10^{-1}	6.5 $\times$




Figure 42. Training and extrapolation MSE vs. noise level, both methods.

KAN-ODE wins on absolute accuracy at every level, but the margin collapses $170\times$ across the sweep (sparse regression degrades $9.0\times$ from $\sigma = 0$ to 0.1 ; KAN-ODE degrades $1540\times$). The sparse regression’s six-term library caps how much noise can move it; the KAN’s 240 free parameters fit the noise directly and the error compounds through the solver. Per-term recovery shows the bilinear cross-term (xy) is recovered to within 5.7% at every noise level – exactly the term a single-layer KAN edge, being strictly univariate, cannot represent directly, since each edge is a function of only one input variable and a layer sums its edges additively. This asymmetry – sparse regression writes xy as a literal library column, while a KAN can only approximate it through composition across a hidden layer – means a term-by-term symbolic comparison would be comparing a coefficient against a distributed representation, so the comparison used here is trajectory reconstruction instead.




Figure 43. Reconstructed trajectories at $\sigma = 0$ and $\sigma = 0.1$: sparse regression loses amplitude under noise; KAN-ODE keeps amplitude but drifts in phase.

8.2 Edge pruning: testing whether a term-by-term comparison is even possible

Part of the original motivation for a smooth, globally-supported basis such as RBF is that a sparse, pruned KAN edge structure could in principle be read back as a closed-form symbolic law, the same interpretability goal sparse regression achieves by construction. The base paper reaches this by adding an L_1 penalty on every spline coefficient to the training loss,

$$\mathcal{L} = \text{MSE}(\mathbf{u}_{\text{true}}, \mathbf{u}_{\text{pred}}) + \lambda \sum_i |c_i|, \quad (21)$$

where c_i ranges over every spline coefficient in Equation 2 and λ trades off trajectory accuracy against how many edges the penalty drives toward zero. Every checkpoint trained for this project used $\lambda = 0$: the plain mean-squared-error term alone, with no sparsity pressure ever added to Equation 21. A KAN-ODE analogue of a sparse coefficient list would instead come from pruning weak edges *post hoc* and reading off what survives. Ranking each edge by mean spline-coefficient magnitude and masking the weakest, applied to the existing (never retrained) noise-sweep checkpoints, removes the two weakest of forty edges and raises the noise-free full-horizon MSE from 8.90×10^{-5} to 3.43 – a factor of 3.9×10^4 , worse than predicting the constant dataset mean (2.91). Every checkpoint was trained with no sparsity-inducing regularization applied at all, so nothing during training ever asked the network to concentrate its fit onto a few edges; the flat plateau at ≈ 3.5 – 3.8 across all pruning levels of 25% or more is the signature of a model where every edge carries part of the fit, confirming that the pruning procedure itself works correctly on a network that has no prunable structure to find.




Figure 44. Full-horizon MSE vs. edge-pruning fraction, all four noise levels.

8.3 Testing whether SIR's fixes transfer to a non-mechanistic epidemic curve

A single outbreak wave that no compartmental ODE generated tests whether the SIR fixes of Section 3 generalize by mechanism rather than by coincidence of the SIR system's own structure. The curve is synthetic but shaped like observed case data: a Gaussian wave centred at day 38 and skewed so that its infected peak falls on day 45, with 1% observation noise and a 7-day moving average, spanning 120 days, of which the first 45 are used for training. The same init-time diagnostic used to root-cause SIR reproduces an almost identical speed-mismatch signature ($18.77\times$ vs. SIR's $19\times$), justifying a time rescaling by a factor of 24 *a priori*, before any training.

Question	SIR's three stability fixes were each derived from a property of the SIR generator itself – do they still work on data no compartmental model produced?
Expected	Only the unit-rescaling fix, which assumes nothing about the data's structure, should transfer; the two structural fixes should not.
Observed	One in three predictions was wrong: the vanishing-field gate transfers and is the best single option, despite being predicted inapplicable.



Table 22. Which of SIR's three fixes transfer to a non-mechanistic curve.

Fix	Precondition	Holds?	Outcome
Time rescaling	None – a unit change	Yes	Removes a $263,049\times$ gradient spike
Conservation via sub-space projection	An exact invariant, $\sum y = c$	No ($\sum y$ spans 0.002–1.54)	Inapplicable; $25\times$ worse fit
Vanishing-field gate at the equilibrium point	$y_d = 0$ is a manifold of equilibria	Yes (unexpectedly)	Best single option, $47\times$ better fit

The vanishing-field gate was *predicted* inapplicable: the infected fraction at $t = 0$ is 0.0027, seemingly too small to avoid freezing the field entirely. A control arm overturns that prediction: gating the learned field by the initial state correctly encodes that the rate of change is proportional to current prevalence, and is the only configuration that reproduces the existence of a peak (day 48 predicted vs. true day 45) from a training window that never observed one.

The extrapolation runaway traces to the train/test split. The default split lands exactly on the infected peak (0% of the decay phase is in the training window), so every model is asked to predict a monotone decay it has never observed. Moving the split just 15 days later collapses the predicted maximum from $6.18\times$ the true peak to 0.964 (true maximum = 1.0), turning over at day 47 against a true day 45 – the runaway is a property of where the window ends, not of the method fitting the data.

Time-scale sweep on the empirical outbreak (2,000-epoch probes)

**Figure 45.** Time-rescaling sweep: accuracy and gradient-spike ratio vs. the rescaling factor.**Figure 46.** Outbreak-curve fit: default split (runaway), the vanishing-field-gate arm, and the fifteen-day-later split control.**Verdict**

Time rescaling transfers unconditionally because it is a unit change with no structural precondition. The other two fixes are structural priors that transfer exactly insofar as their structure is actually present in the new data: the conservation projection's precondition is confirmable in advance by directly measuring the compartment sum, while the vanishing-field gate's precondition turned out to require running the control arm to settle — argument alone predicted the wrong outcome. That arm cost 27 minutes and overturned the prediction.

9 Validation: Epoch-Budget Sensitivity and Implementation Audit

Every result in Sections 2–8 was measured at a fixed 10^4 -epoch budget, a tenth of the base paper’s, for the cost reasons of Section 2.3. This stage asks whether that budget is long enough for the *rankings* to be trusted, and whether any gap to the paper is due to the budget or to an implementation difference.

Question	Do the conclusions drawn at 10^4 epochs survive five times the training budget?
Expected	Absolute errors fall; orderings between configurations hold, since every 10^4 -epoch model had already reached a training loss near 10^{-4} .
Observed	Three of four conclusions hold. The fourth, the one the paper headlines, reverses: at 5×10^4 epochs the parameter-matched MLP extrapolates $6.8\times$ <i>better</i> than KAN-ODE.

9.1 Design

Five configurations were chosen to test the specific conclusions most exposed to the budget: Euler (the one solver short of 10^{-4} at 10^4 epochs), Tsit5/RBF (the reference), B-spline (the near-tie with RBF), and both MLP baselines. Each was retrained from scratch at 5×10^4 epochs with the recipe of Table 1 unchanged, on cloud CPUs. The B-spline run was cut to 2.5×10^4 epochs because its short calibration run predicted about 35 hours for the full budget; the real run turned out $1.63\times$ faster than that calibration (1.544 against 2.52 s/epoch), a reminder that a five-epoch sample is dominated by start-up cost. The batch used 23.2 serial CPU-hours against a planned 17.0, and 10.7 hours of wall-clock with the five jobs in parallel. Before analysis every run’s stored metrics were checked against the epoch count of its training log, and every checkpoint was reloaded and re-scored; the rollouts reproduce the stored extrapolation MSE to within 0.6%.


A sixth configuration was added later, as a separate follow-up: the implementation audit (Section 9.3) turned up the base paper’s own learning rate and initialization scale for its MLP baseline, and both budgets were rerun with the identical $[2, 50, 2]$ tanh architecture under those settings instead of this project’s default. It is reported alongside the original five throughout this section.

9.2 Results


Table 23. The five originally planned configurations at 10^4 epochs (Section 3) and at the extended budget, plus a sixth added after auditing the paper’s own optimizer settings for its MLP baseline (Section 9.3). “Clip” is the fraction of optimizer steps on which gradient clipping at norm 1.0 engaged; $\|\nabla\|_{\text{med}}$ is the median pre-clip gradient norm over the extended run. The sixth configuration ran with no clipping by design (see the paragraph following this table), so those two columns are blank throughout.

Configuration	Epochs	Training MSE	Extrap. MSE	Extrap. R^2	$\hat{L}$	Clip	$\ \nabla\ _{\text{med}}$
Euler / RBF	10k	1.38×10^{-4}	2.51×10^{-3}	0.99914	8.55		
	50k	8.85×10^{-6}	1.93×10^{-4}	0.99993	7.72	4.7%	0.147
Tsit5 / RBF	10k	8.83×10^{-5}	8.92×10^{-5}	0.99997	6.91		
	50k	6.59×10^{-6}	4.02×10^{-5}	0.999986	7.75	4.3%	0.169
Tsit5 / B-spline	10k	7.20×10^{-5}	8.38×10^{-5}	0.99997	8.44		
	25k	9.65×10^{-6}	4.26×10^{-5}	0.999985	7.05	2.2%	0.033
MLP-ODE, SiLU	10k	9.50×10^{-5}	4.77×10^{-4}	0.99984	5.07		
	50k	1.82×10^{-6}	5.90×10^{-6}	0.999998	5.65	11.1%	0.162
MLP-ODE, tanh (paper, default lr/init)	10k	1.08×10^0	1.69×10^0	0.4215	38.2		
	50k	1.01×10^0	4.32×10^0	−0.4841	104.8	99.3%	19.8
MLP-ODE, tanh (paper’s own lr/init)	10k	1.11×10^{-3}	9.65×10^{-3}	0.99669	9.23		
	50k	5.43×10^{-5}	4.31×10^{-4}	0.99985	9.67	–	–


Every working configuration improves its training MSE by 7.5–52 $\times$ (Table 23); what matters is which order-




(a) KAN-ODE, Tsit5 / RBF, 50k




(b) MLP-ODE, SiLU, 50k




(c) KAN-ODE, Euler / RBF, 50k



(d) KAN-ODE, Tsit5 / B-spline, 25k



(e) MLP-ODE, tanh (paper's default lr/init), 50k



(f) MLP-ODE, tanh (paper's own lr/init), 50k

Figure 47. Training and test loss of the six extended runs, as saved by each run. Training loss is logged every epoch; the test loss (full horizon $[0, 14]$) every ten epochs, and it spikes by up to two orders of magnitude throughout, which is why the comparison in the text uses a 500-epoch rolling median rather than raw values. Panel (f) shows the identical $[2, 50, 2]$ tanh architecture of panel (e), retrained with only the learning rate and initialization scale changed to the values in the paper's own published code.

ings survive. Table 24 states the verdict for each conclusion the earlier sections draw.

Table 24. Which 10^4 -epoch conclusions survive the extended budget.

Conclusion at 10^4 epochs	At 10^4	At 5×10^4	Verdict
RBF and B-spline are statistically tied	6.1% apart	5.6% apart (25k)	Holds. The basis choice can be made on cost.
Euler is the worst solver	$28.1 \times$ Tsit5	$4.8 \times$ Tsit5	Holds , with a much smaller margin: part of the gap was under-training.
The paper-spec MLP fails because of its architecture, not its budget	train 1.08, $R^2 = 0.42$	train 1.01, $R^2 = -0.48$	Budget part holds; architecture part reversed. More epochs alone do not help, but the identical architecture trains under the paper’s own learning rate and initialization (below): $R^2 = 0.9998$ at this budget.
KAN-ODE extrapolates better than a parameter-matched MLP	KAN $5.35 \times$ better	MLP $6.82 \times$ better	Reverses. The advantage is a budget effect.

The paper-spec MLP is stuck, not slow. With five times the budget its training loss moves by 7% ($1.083 \rightarrow 1.011$), its extrapolation error rises $2.6 \times$, and R^2 turns negative. The optimizer spends the run in a pathological regime: clipping engages on 99.3% of steps against 2–11% for every other configuration, the median pre-clip gradient norm is 19.8 against 0.03–0.17, and the Lipschitz estimate nearly triples to 104.8, an order of magnitude above every other model. This settles half of the question left open in Section 2.4: whatever is wrong here, it is a property of this training run, not of the epoch count, since five times the budget only entrenches it. It is not, however, a property of the architecture itself – the next paragraph shows the identical network training cleanly once its optimizer settings are corrected.

Resolving the audit’s open optimizer row: the paper’s own learning rate and initialization. The base paper’s own published code (not distributed with this project, but readable in its public repository) sets two things differently for its MLP baseline than this project’s default: a learning rate of 10^{-2} (five times this project’s 2×10^{-3}) and Glorot initial weights divided by 10^5 , near zero rather than full-scale. Holding architecture, activation, dataset, solver and seed fixed and changing only those two settings, the identical $[2, 50, 2]$ tanh network trains cleanly at both budgets (Table 23, row six; panel f of Figure 47): at 10^4 epochs, training MSE falls from 1.08 to 1.11×10^{-3} and extrapolation MSE from 1.69 to 9.65×10^{-3} ($R^2 : 0.42 \rightarrow 0.997$); at 5×10^4 epochs training MSE reaches 5.43×10^{-5} and extrapolation MSE 4.31×10^{-4} ($R^2 = 0.99985$), against 1.01 and 4.32 under the default settings at the same budget. Neither run needed gradient clipping or encountered a single non-finite gradient step. The conclusion Table 4 drew does not survive this: a change of only learning rate and initialization scale, with the exact literal architecture and activation that table showed failing, is enough on its own. What that factorial actually measured was sensitivity to activation *at this project’s optimizer settings*, not an architecture-level incompatibility with the system.

Against the other two architectures at 5×10^4 epochs, the corrected tanh MLP is still not competitive: its extrapolation MSE (4.31×10^{-4}) is $10.7 \times$ KAN-ODE’s and $73 \times$ the SiLU MLP’s, and its learned equilibrium sits 1.72 from the truth, the furthest of the three (Table 10). Its one advantage is conservation: a first-integral drift of 0.06% is the smallest of the three (KAN 0.08%, SiLU-MLP 0.21%), though its on-orbit field error (2.7%) is the largest. It also falls short of the paper’s own reported target: the paper states its MLP reaches a training loss of 3.0×10^{-5} within 10^5 epochs, and this run’s training loss crosses 10^{-4} at epoch 39,695 but had not reached 3.0×10^{-5} by epoch 50,000 – on a trend consistent with, but not confirming, the paper’s own budget for this baseline.

The KAN advantage over the MLP reverses. At 10^4 epochs, from training losses 7% apart, the KAN-ODE extrapolates $5.35 \times$ better than the SiLU MLP, the result that reproduces the paper’s accuracy claim. At 5×10^4 epochs the MLP’s extrapolation MSE is 5.90×10^{-6} and the KAN’s 4.02×10^{-5} : the MLP is $6.82 \times$ better, far outside the single-seed noise that separates RBF from B-spline. The checkpoint analyses of Section 3.5 agree on every measure: at 50k the MLP has the smaller on-orbit field error (0.6% against 1.1%), the smaller first-integral drift (0.4% against 0.9%), an equilibrium closer to the truth (0.68 against 1.65 away), and a lower

far-horizon error (maximum 2.4×10^{-2} against 6.5×10^{-2} on (14,28]). The crossover is late. Smoothing the logged full-horizon MSE of Figure 47 with a 500-epoch rolling geometric median, the two models are level at 2.5×10^4 epochs (1.7×10^{-4} against 2.0×10^{-4}), and the MLP stays ahead for good only from about epoch 48,500. The MLP is also the faster learner at every threshold: in the extended runs it reaches 10^{-4} training MSE at epoch 5225 against the KAN’s 5695, and 10^{-5} at 13,815 against 31,412.

The honest reading is that KAN-ODE has a strong *early* inductive bias for this smooth periodic problem. Within a fixed, moderate budget it produces the better extrapolator, but it is not the better function class once both models are trained to convergence. Neither model learns the conservative structure of the system (Figure 16); the MLP simply fits the training orbit’s phase more precisely in the end.

9.3 Implementation audit against the base paper

The second validation task compares this project’s recipe with the paper’s stated methodology, field by field, so that no difference is silently absorbed into “the paper trained longer”. The paper itself is not in the project repository; Table 25 uses the settings recorded in the project’s own planning documents and code comments, and marks as open every field those do not determine. One such field, the MLP baseline’s optimizer settings, was later resolved by reading the paper’s own published code directly, which is how the follow-up above was found.

Table 25. Field-by-field comparison with the base paper [6].

Field	Base paper	This project	Verdict
KAN architecture	[2, 10, 2], $G = 5$, 240 parameters	identical	Matches.
Edge basis	Gaussian RBF, $\exp(-r^2/2h^2)$	identical, after fixing a missing factor $\frac{1}{2}$	Matches. The fix predates every reported run.
MLP baseline	[2, 50, 2], tanh, 252 parameters	identical, plus a working [2, 14, 8, 8, 2] SiLU variant	Matches literally, and now trains under the paper’s own optimizer settings (below).
Integrator	adaptive Tsit5 (Julia)	fixed-step Tsit5, $h = 0.05$	Explained: the fixed step’s own error is 4.7×10^{-14} in MSE (Table 7).
Gradients	continuous adjoint	direct autograd	Explained: differ in memory and time, not in the loss minimised (Section 4).
Framework	compiled Julia / SciML	interpreted PyTorch, CPU	Explained: wall-clock is not comparable across the two (Section 2.3).
Epoch budget	10^4 (KAN), 10^5 (MLP)	10^4 , and 5×10^4 here	Explained: this section.
Optimizer, learning rate (MLP baseline)	Adam, 10^{-2} , Glorot init $\div 10^5$, found in the paper’s own published code	Adam, 2×10^{-3} by default, Glorot gain 1.0, clip 1.0	Resolved: this mismatch, not the architecture, caused the earlier training failure (below). The paper’s own KAN learning rate remains unconfirmed.
Loss normalisation	not recorded	mean over points and components	Open; affects absolute loss values only.

The paper reports its KAN-ODE reaching a training loss of 2.6×10^{-5} in 10^4 epochs and its MLP needing 10^5 epochs for 3.0×10^{-5} . The extended runs locate both targets directly. This project’s KAN-ODE reaches 2.6×10^{-5} at epoch 13,632, within $1.4\times$ of the paper’s KAN. Its SiLU MLP reaches 3.0×10^{-5} at epoch 8,585, nearly $12\times$ sooner than the paper’s MLP. The literal [2, 50, 2] tanh MLP never gets below 1.0 under this project’s default optimizer settings, but under the paper’s own learning rate and initialization (above) it crosses 10^{-4} at epoch 39,695 and had not reached 3.0×10^{-5} by epoch 50,000 – a trend consistent with, but short of confirming, the paper’s own 10^5 -epoch target. Allowing for the open loss-normalisation row, the KAN side of the paper’s efficiency claim is reproduced; its tenfold advantage over the literal MLP baseline turns out to depend on that baseline’s optimizer settings as much as on its architecture, since the identical network trains once those settings are corrected.

Validation outcome

The 10^4 -epoch budget gives trustworthy *rankings* for the solver and basis ablations: the RBF/B-spline tie and Euler’s last place both survive five times the budget. It does not for the KAN-versus-MLP comparison, which reverses: the paper’s extrapolation advantage is an early-training effect under this protocol, and its speed advantage over the literal MLP baseline depended on that baseline’s default optimizer settings, not on its architecture – corrected, the identical network trains, though it still does not close the gap with either KAN-ODE or the SiLU MLP within this budget. Every result is single-seed.

10 Discussion

Eight patterns recur across studies that were designed independently and address unrelated questions. Their recurrence is itself informative about how this class of model should be studied going forward.

A trained Neural ODE models its training orbit and its solver, not the vector field. The numerical analysis of Section 3.5 found this in three independent ways. The Euler-trained network is an accurate model of the inverse modified equation, and no model transfers across solver families (Figure 9). Every trained model, KAN or MLP, at either budget, has an attracting limit cycle through the training data around a spurious unstable equilibrium 1.2–1.7 away from the true one, where Lotka-Volterra has a neutral centre. And the field is accurate only in a thin tube around the observed orbit (Figure 17). Low trajectory error on the training orbit’s continuation is therefore compatible with a qualitatively wrong phase portrait. Checks that probe the field directly, such as locating its equilibria, integrating from new initial conditions or evaluating a known invariant, are cheap once a checkpoint exists, and they should be part of the protocol whenever a learned field is to be read as physics.

Architecture comparisons are budget-dependent. The KAN-versus-MLP ranking reversed between 10^4 and 5×10^4 epochs (Section 9), while the solver and basis rankings did not. The difference is instructive: the solver and basis comparisons are between models of the same function class, and they were decided by truncation error or locality, properties that do not change with training time. The architecture comparison is between two function classes that converge at different rates, and a fixed budget measures the rates as much as the limits. A claim that one architecture “generalises better” should state the budget, or better, report the error against training time for both.

A reasoned hypothesis still requires a control arm. Three separate studies advanced a specific, mechanistically plausible hypothesis and then tested it directly: gradient-norm roughness as the cause of Euler’s convergence floor (Section 5), SiLU’s representational limit as the cause of the pendulum’s early failure (Section 3), and the predicted inapplicability of the vanishing-field gate to non-mechanistic data (Section 8). Each hypothesis had a clear argument behind it before any run was launched, and each was overturned by the run itself. The result reported in every one of these three cases is the outcome of the control arm, and the original argument survives only as the motivation for running it. This is not a coincidence of three unlucky guesses; it reflects the difficulty of reasoning about a system where the integrator, the optimizer, and the learned function interact inside the same backpropagation graph. The standard practice of ablating one variable while holding others fixed rests on an assumption of near-independence that does not hold here.

A single seed and a single train/test split account for results that first looked like properties of the solver or the method. The stiffness-map failures at $\mu \in \{0.5, 1.0, 2.0\}$ (Section 7) and the epidemic-fit extrapolation runaway (Section 8) both appeared, on first inspection, to indict a solver or a modeling approach. Retraining with additional seeds, in the stiffness-map case, and shifting the train/test split by fifteen days, in the epidemic case, resolved both into an artifact of one specific run’s random initialization or one specific window boundary. Neither observation survives as a claim about solver order or about KAN-ODE’s capacity once the seed or the split is varied. Every quantitative result in this report remains $N = 1$; the stiffness study’s small extension to $N = 3$ was enough to dissolve a headline finding, which underscores how much of the remaining project rests on the same single-seed foundation.

Measuring a quantity before training is more informative than tuning after a failed run. The SIR divergence (Section 3), the outbreak-curve fit (Section 8), and the stiffness map’s spurious equilibrium (Section 7) were each root-caused by evaluating the model’s own initial field, or its trained field at a candidate equilibrium, and comparing that number against the system’s true scale. In every case this produced a specific, falsifiable prediction – a rescaling factor, a fix that should or should not transfer, an equilibrium location – that a subsequent short run confirmed or refuted directly, in contrast to the open-ended search a hyperparameter sweep would have required.

Training accuracy and generalization quality are decoupled. The hybrid basis matches pure RBF’s training loss to 1.7% at 10,000 epochs yet collapses in extrapolation ($R^2 = -2.17$ vs. $+0.647$, Section 6). The stiffness map’s trapped models (Section 7) reach training losses of 1.8×10^{-2} – qualitatively acceptable on a casual inspection – while their learned field has a large nonzero derivative at the true equilibrium. And the epidemic-curve model (Section 8) produces a plausible training fit while predicting an infected peak six times the physical maximum. In all three cases, a reader who judged the model solely by training-window MSE would reach the wrong conclusion. The evaluation protocol established in Section 2 – scoring strictly on extrapolation, never on training loss – was adopted precisely because this decoupling was anticipated, but these studies go further: they show that a model can produce numerically indistinguishable training fits from configurations that differ by two orders of magnitude in extrapolation quality. The number to watch is always the one computed outside the training window, and no training-window metric tested in this project reliably predicts it.

Cost and accuracy form a clear Pareto frontier. Across the solver ablation (Section 3, Figure 10), the basis ablation, and the hybrid-basis study (Section 6), a consistent pattern emerges: the cheapest methods that work (Midpoint at $1/3$ the function evaluations of Tsit5; RBF at $1/3.9$ the wall-clock of B-spline) match the accuracy of their more expensive alternatives to within single-seed noise. Adding cost – more function evaluations per step, more basis evaluations per edge, a second basis evaluated in parallel – does not buy a measurable accuracy improvement on any system tested. The hybrid-basis result makes this particularly sharp: the blend pays roughly the sum of both bases’ per-epoch compute for two extra parameters and never recovers the cost in accuracy.

These KAN edges do not produce interpretable symbolic structure without explicit encouragement. The pruning experiment (Section 8) found that removing the two weakest of forty edges degrades a noise-free model to worse than a constant predictor – a factor of 3.9×10^4 . Every edge carries part of the fit, the magnitude ranking is close to arbitrary, and no edge can be identified with a specific term in the governing equation. This is a direct consequence of training without sparsity-inducing regularization, and it means the interpretability advantage that motivated the original choice of a smooth, globally-supported basis function remains a hypothesis rather than a demonstrated property of the trained models in this project. A future retraining with explicit edge-level regularization would be needed to test whether sparse, readable edge functions emerge when the training loss actively asks for them.

11 Conclusion

This project reproduced a published KAN-ODE recipe on the Lotka-Volterra predator-prey system with every integrator and basis implemented from its defining formula, verified those implementations numerically, extended the recipe to two additional dynamical systems (a damped pendulum and an SIR epidemic model), diagnosed and fixed two genuine cross-domain training failures, ran five further studies that each isolated a specific question about solver behavior, gradient dynamics, basis-function design, stiffness, and comparison against a fundamentally different modeling paradigm, and validated the budget-sensitive conclusions at up to five times the training budget.

The principal quantitative findings are:

- All six integrators satisfy their Runge-Kutta order conditions to round-off (37 rooted trees checked) and show their design order on the true field; Tsit5’s principal error norm is $2.9\times$ smaller than DOPRI5’s. Above order two the solver does not change accuracy at the training step, and Midpoint matches Tsit5 within 0.3% at 29% of the wall-clock.

- The network learns the field its training solver needs: the Euler model matches the inverse modified equation $\mathbf{f} + \frac{h}{2}J\mathbf{f}$ (cosine 0.92, magnitude ratio 1.04), and models lose one to four orders of magnitude when integrated with a solver from another family.
- Every trained model, KAN or MLP, is an attracting limit-cycle oscillator around a spurious unstable equilibrium, where Lotka-Volterra has a neutral centre; the field is accurate only in a thin tube around the training orbit.
- B-spline and RBF are tied on accuracy at both budgets; locality, not conditioning, separates the bases, and the global polynomials extrapolate $4\text{--}5\times$ worse. Extrapolation MSE scales as $12.9\sigma^{2.02}$ with observation noise.
- At 10^4 epochs KAN-ODE extrapolates $5.4\times$ better than a parameter-matched MLP, a gap explained by a $48\times$ smaller period error rather than a more accurate field; the KAN advantage widens to $40.6\times$ at twice the horizon (Table 15). At 5×10^4 epochs the comparison reverses and the MLP is $6.8\times$ better. The paper’s literal tanh MLP fails under this project’s default optimizer settings, but trains under the paper’s own learning rate and near-zero initialization scale, reaching $R^2 = 0.9998$ at 5×10^4 epochs, still $73\times$ behind the SiLU MLP on extrapolation MSE; the paper’s tenfold speed claim compared against the uncorrected baseline, not an architectural limitation.
- SIR’s training failure traced to a $19\times$ unit-scale mismatch between the network’s initial field and the true dynamics, fixable by a change of time units and a conservation projection – two interventions that assume nothing about the system beyond what its own equations state and that together remove the divergence. Adding a vanishing-field gate moved extrapolation R^2 from -20.97 to $+0.9981$.
- Direct autograd is the correct default for this project’s scale: adjoint sensitivity already uses less memory at every tested trajectory length but costs a consistent $1.9\text{--}2.1\times$ wall-clock overhead.
- Gradient-norm roughness does not explain Forward Euler’s convergence floor (Spearman $\rho = -0.088$, $p = 0.87$); the modified-equation analysis identifies the mechanism as a systematic, solver-specific bias in the learned field.
- A learnable blend of B-spline and RBF basis functions never outperforms the better pure basis on either system tested.
- Stiffness-map failures at transitional damping ratios trace to random initialization rather than solver order or numerical stiffness, confirmed mechanistically by a genuine spurious equilibrium in the learned vector field.
- KAN-ODE is more accurate than sparse symbolic regression at every noise level tested, but sparse regression is far more robust ($9\times$ vs. $1540\times$ degradation from clean to $\sigma = 0.1$). Two of the SIR stability fixes transfer to a non-mechanistic outbreak curve; the one that transfers best was predicted inapplicable.

Every quantitative claim in this report traces to a saved metrics file, a re-scored saved checkpoint, or a regenerable figure, and every study that tested a specific mechanistic hypothesis reports the outcome of a direct control arm rather than the hypothesis alone.

The validation stage settles the budget question for the ablations and overturns it for the architecture comparison, so the principal remaining limitation is statistical: every configuration is a single seed. The stiffness study’s experience, where extending from one seed to three dissolved a headline finding, is the reason multi-seed replication of the KAN-versus-MLP crossover is the most valuable next experiment. Two cheaper follow-ups come directly out of the numerical analysis: an affine input normalisation that uses the whole first-layer grid on positive-valued systems, and training from several initial conditions, the most direct way to ask a Neural ODE for the vector field rather than a single orbit.

References

- [1] Blealtan. efficient-kan, 2024. <https://github.com/Blealtan/efficient-kan>.
- [2] Steven L. Brunton, Joshua L. Proctor, and J. Nathan Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *Proceedings of the National Academy of Sciences*, 113(15):3932–3937, 2016.
- [3] Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David Duvenaud. Neural ordinary differential equations. *Advances in Neural Information Processing Systems*, 31, 2018.
- [4] Germund Dahlquist. A special stability problem for linear multistep methods. *BIT Numerical Mathematics*, 3(1):27–43, 1963.
- [5] Ernst Hairer, Syvert P. Nørsett, and Gerhard Wanner. *Solving Ordinary Differential Equations I: Nonstiff Problems*. Springer, Berlin, 2 edition, 1993.
- [6] Benjamin C. Koenig, Suyong Kim, and Sili Deng. KAN-ODEs: Kolmogorov–Arnold network ordinary differential equations for learning dynamical systems and hidden physics. *Computer Methods in Applied Mechanics and Engineering*, 432:117397, December 2024. ISSN 0045-7825. doi: 10.1016/j.cma.2024.117397. URL <http://dx.doi.org/10.1016/j.cma.2024.117397>.
- [7] Hao Li, Zheng Xu, Gavin Taylor, Christoph Studer, and Tom Goldstein. Visualizing the loss landscape of neural nets. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- [8] Ziming Liu, Yixuan Wang, Sachin Vaidya, Fabian Ruehle, James Halverson, Marin Soljačić, Thomas Y. Hou, and Max Tegmark. KAN: Kolmogorov-arnold networks. *arXiv preprint arXiv:2404.19756*, 2024.
- [9] Ch. Tsitouras. Runge–kutta pairs of order 5(4) satisfying only the first column simplifying assumption. *Computers & Mathematics with Applications*, 62(2):770–775, 2011.